


Which data categories help?

Nine matched 4,000-row training mixes · 12,800 examples seen each · one seed · same 800-row validation set

All mixes used the selected Qwen training settings. Training categories excluded from a mix were replaced by other categories.



Dashed line = control. The paper-share differences are small; leave-one-out differences are larger. This is a single-seed development experiment, not a held-out test or proof of category causality.