



Did the source-balanced retrain fix false alarms?

Each Qwen model uses a 5% false-alarm calibration on the same independent human set.

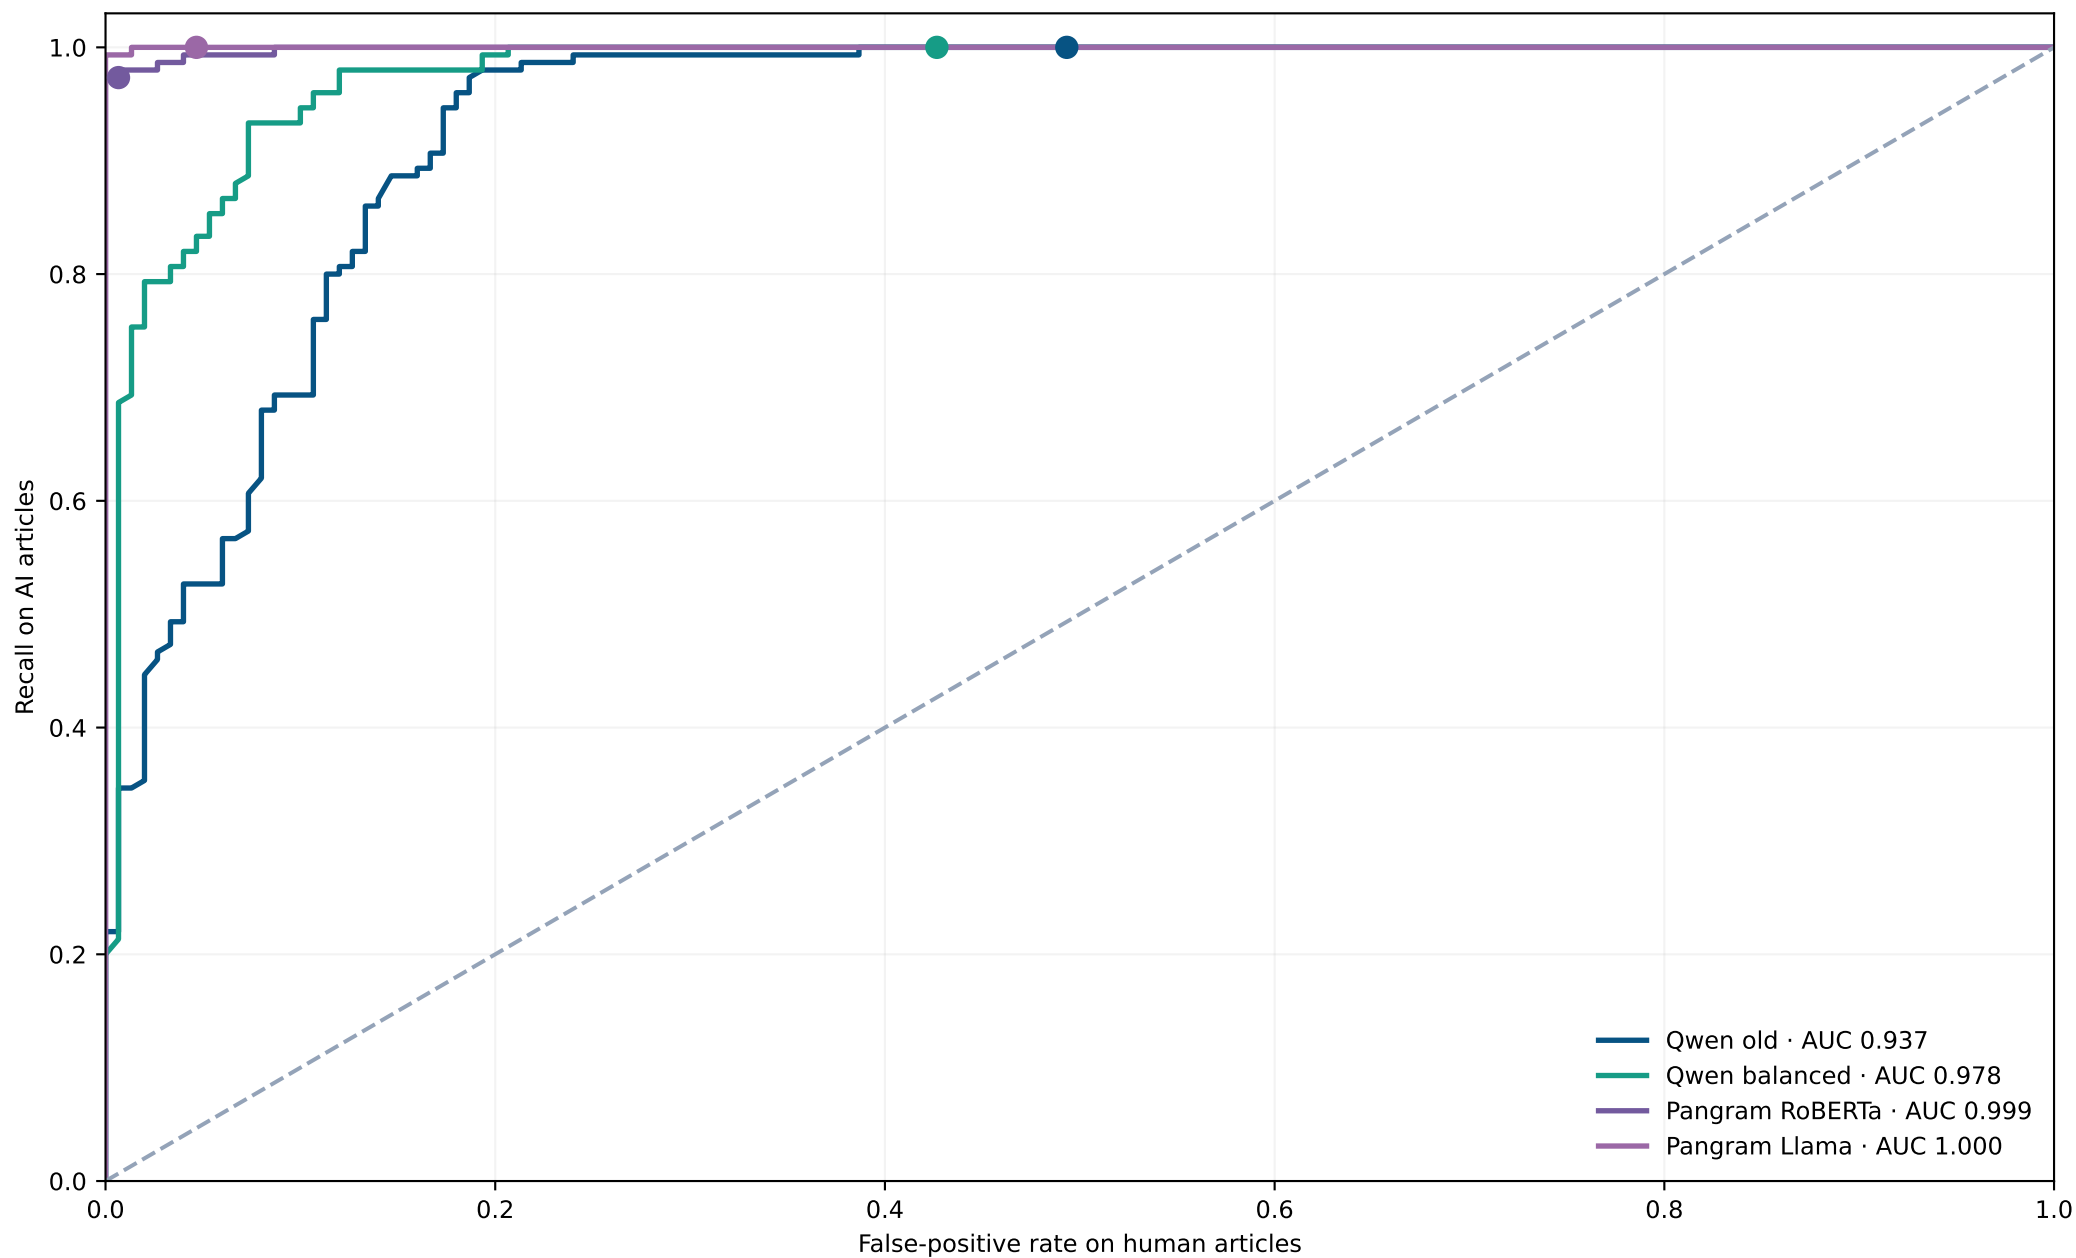
Qwen old Qwen balanced Pangram RoBERTa Pangram Llama



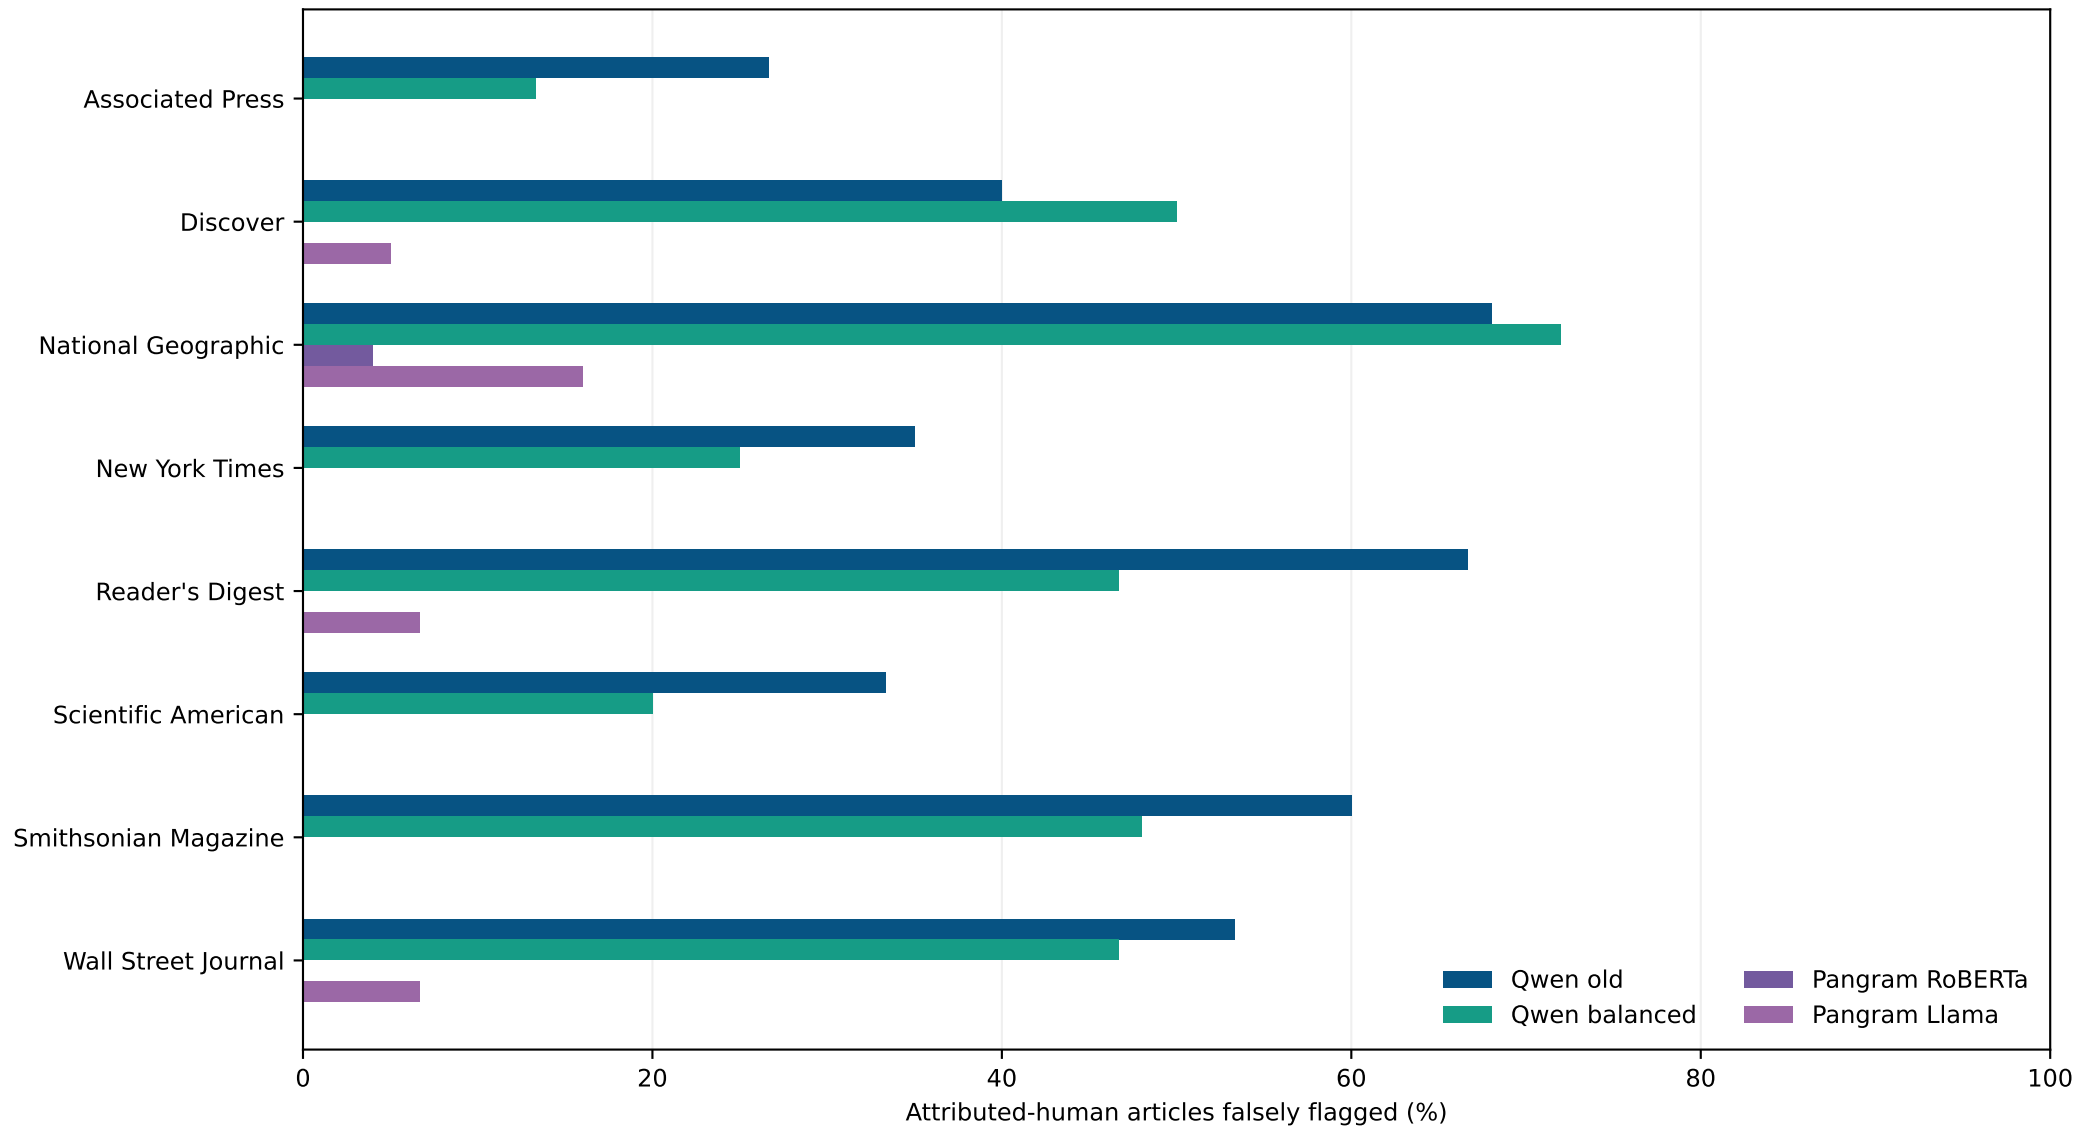
Mixed documents: AI found versus human text falsely marked



Published article ROC: ranking quality across thresholds



Human article false alarms by publication



The external articles remain a frozen stress set. Their named authors do not independently verify AI-free writing workflows.