



Baseline performance | 1,000 balanced test passages | source-aware pyramid v1

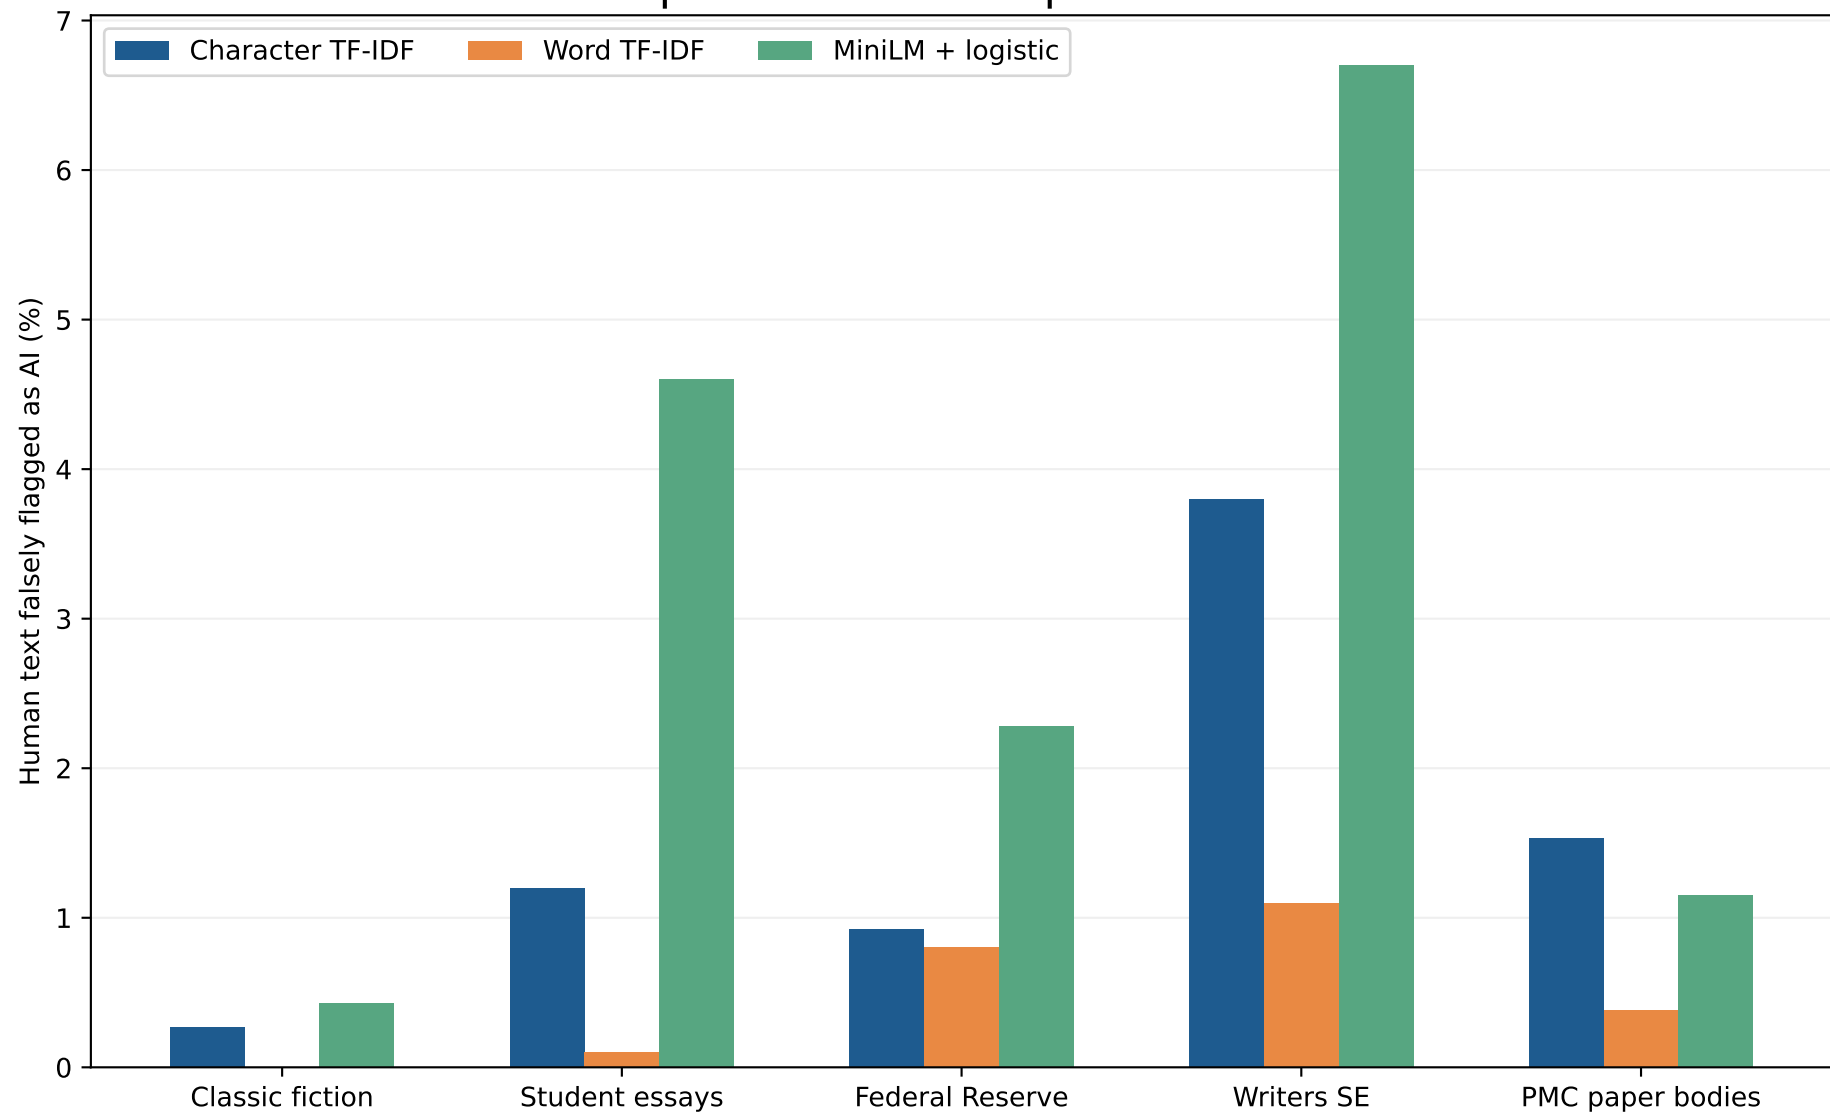
Diverse test: AUROC by writing domain



AI recall at each model's 2% validation human-FPR threshold



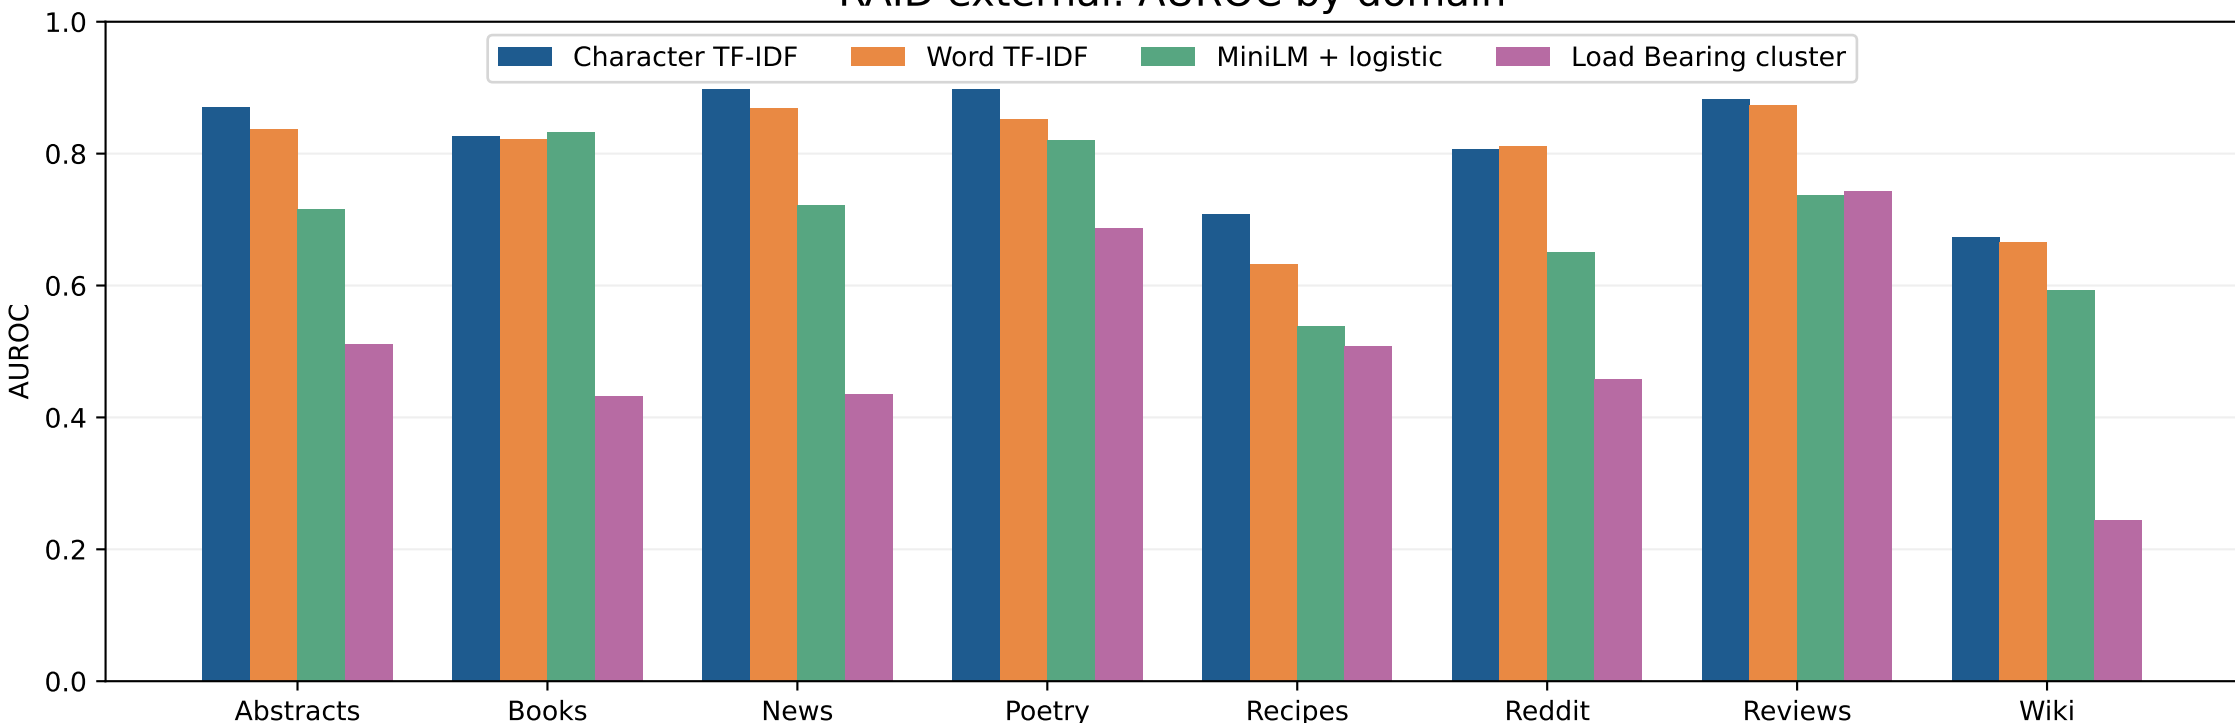
Baseline false positives on independent human sources



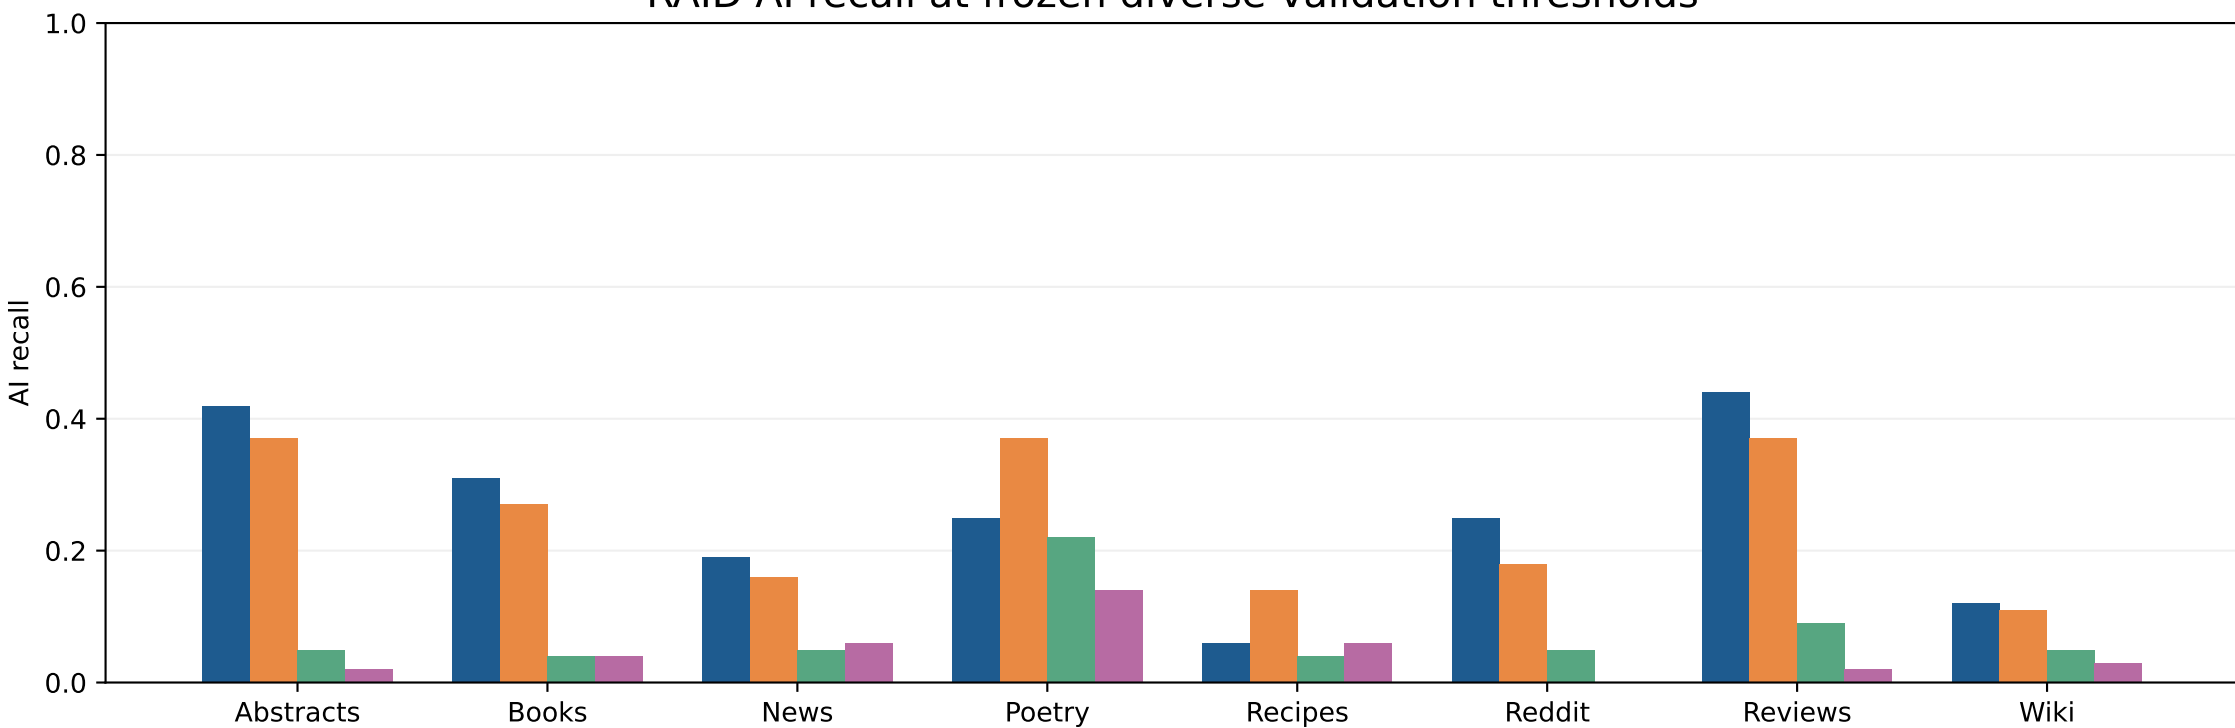
Thresholds calibrated to $\leq 2\%$ human FPR on diverse validation.
Writers Stack Exchange authorship is inferred from dated posts, not individually verified.

Independent source-family test | 1,600 balanced passages | 11 generators

RAID external: AUROC by domain



RAID AI recall at frozen diverse-validation thresholds



Baseline summary

Model	Mixed AUROC	Mixed FPR	Mixed AI recall	GPT-4 OOD AUROC	Paraphrase AUROC	Books human FPR
Character TF-IDF	0.882	1.8%	46.0%	0.889	0.765	0.3%
Word TF-IDF	0.862	1.0%	41.2%	0.893	0.762	0.0%
MiniLM + logistic	0.696	1.6%	9.6%	0.742	0.691	0.4%
Load Bearing cluster	0.481	3.6%	3.6%	0.462	0.385	—

All thresholds selected on the diverse validation split only.
MAGE GPT-4 and paraphrase sets share the MAGE source family with training.
Standard Ebooks is human-only, so AUROC is undefined.