



Diverse test | 1,000 balanced passages | 7 baselines/models

AUROC by domain

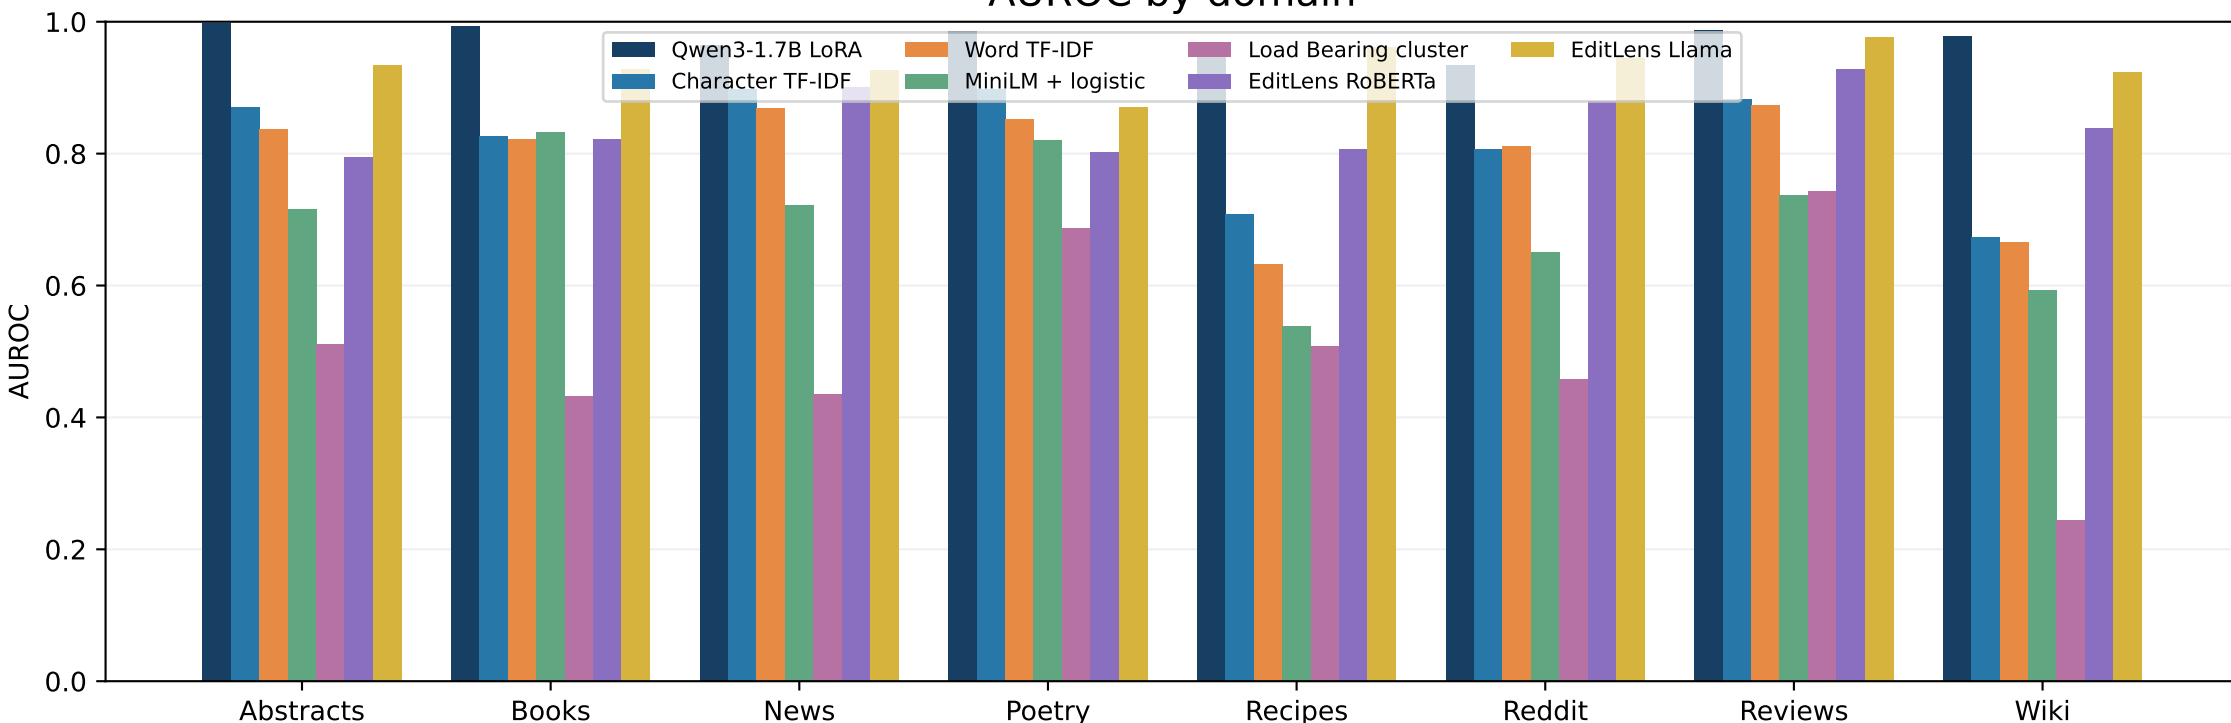


AI recall at each model's validation-calibrated threshold

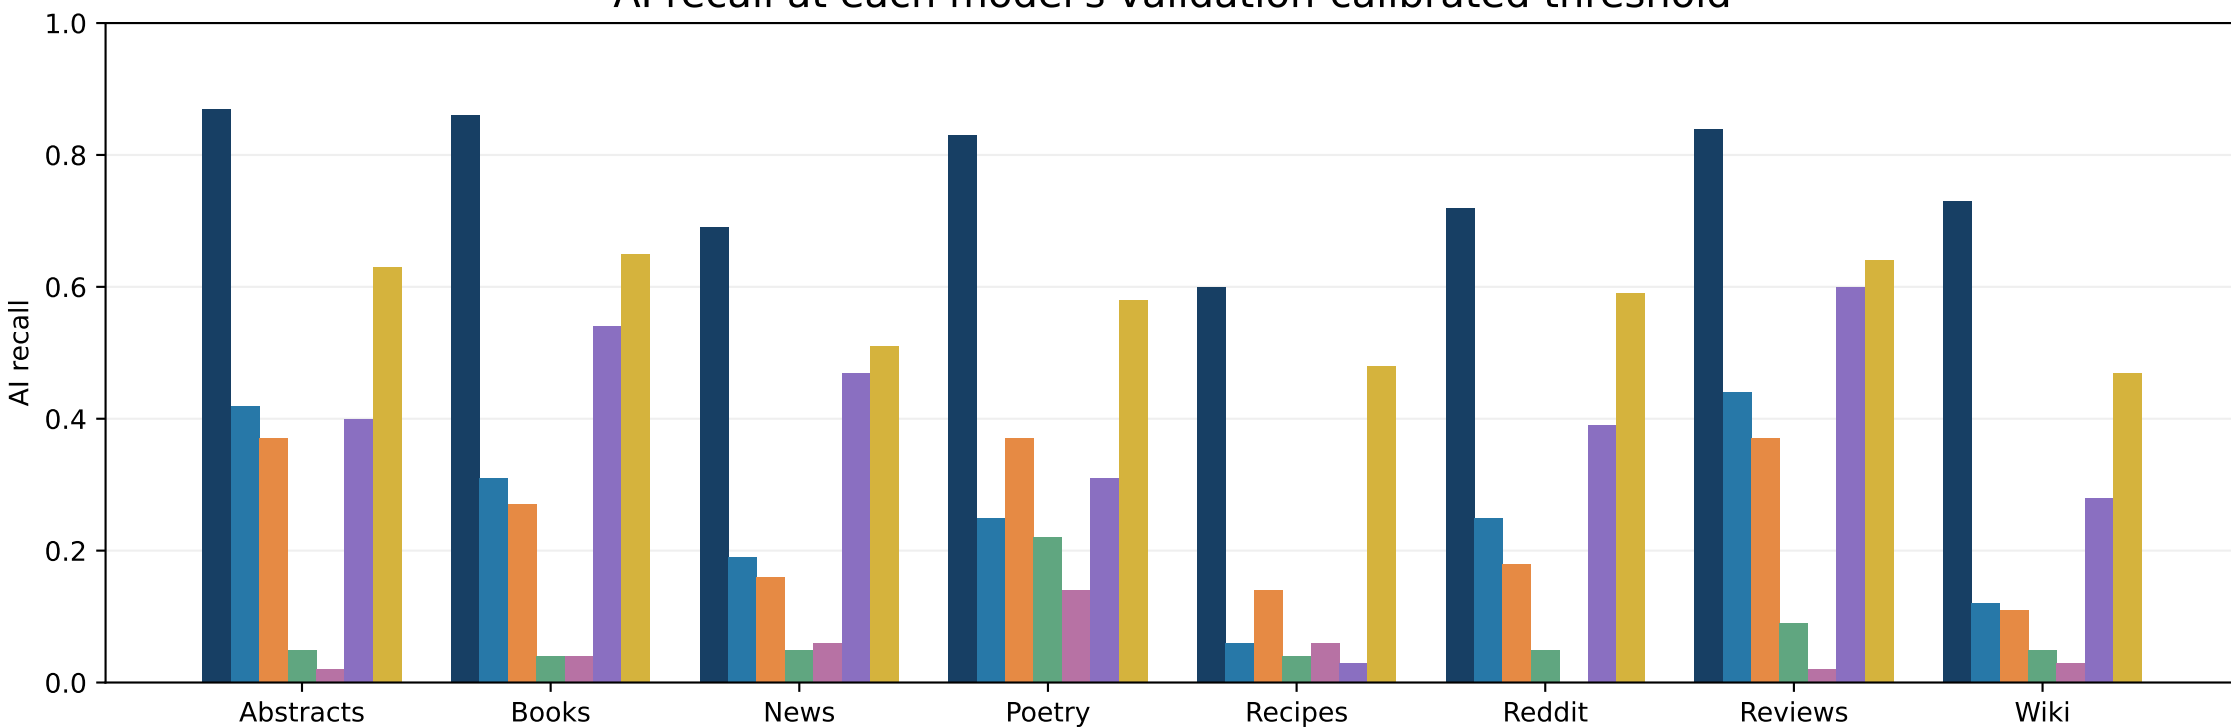


RAID external | 1,600 balanced passages | eight domains

AUROC by domain



AI recall at each model's validation-calibrated threshold

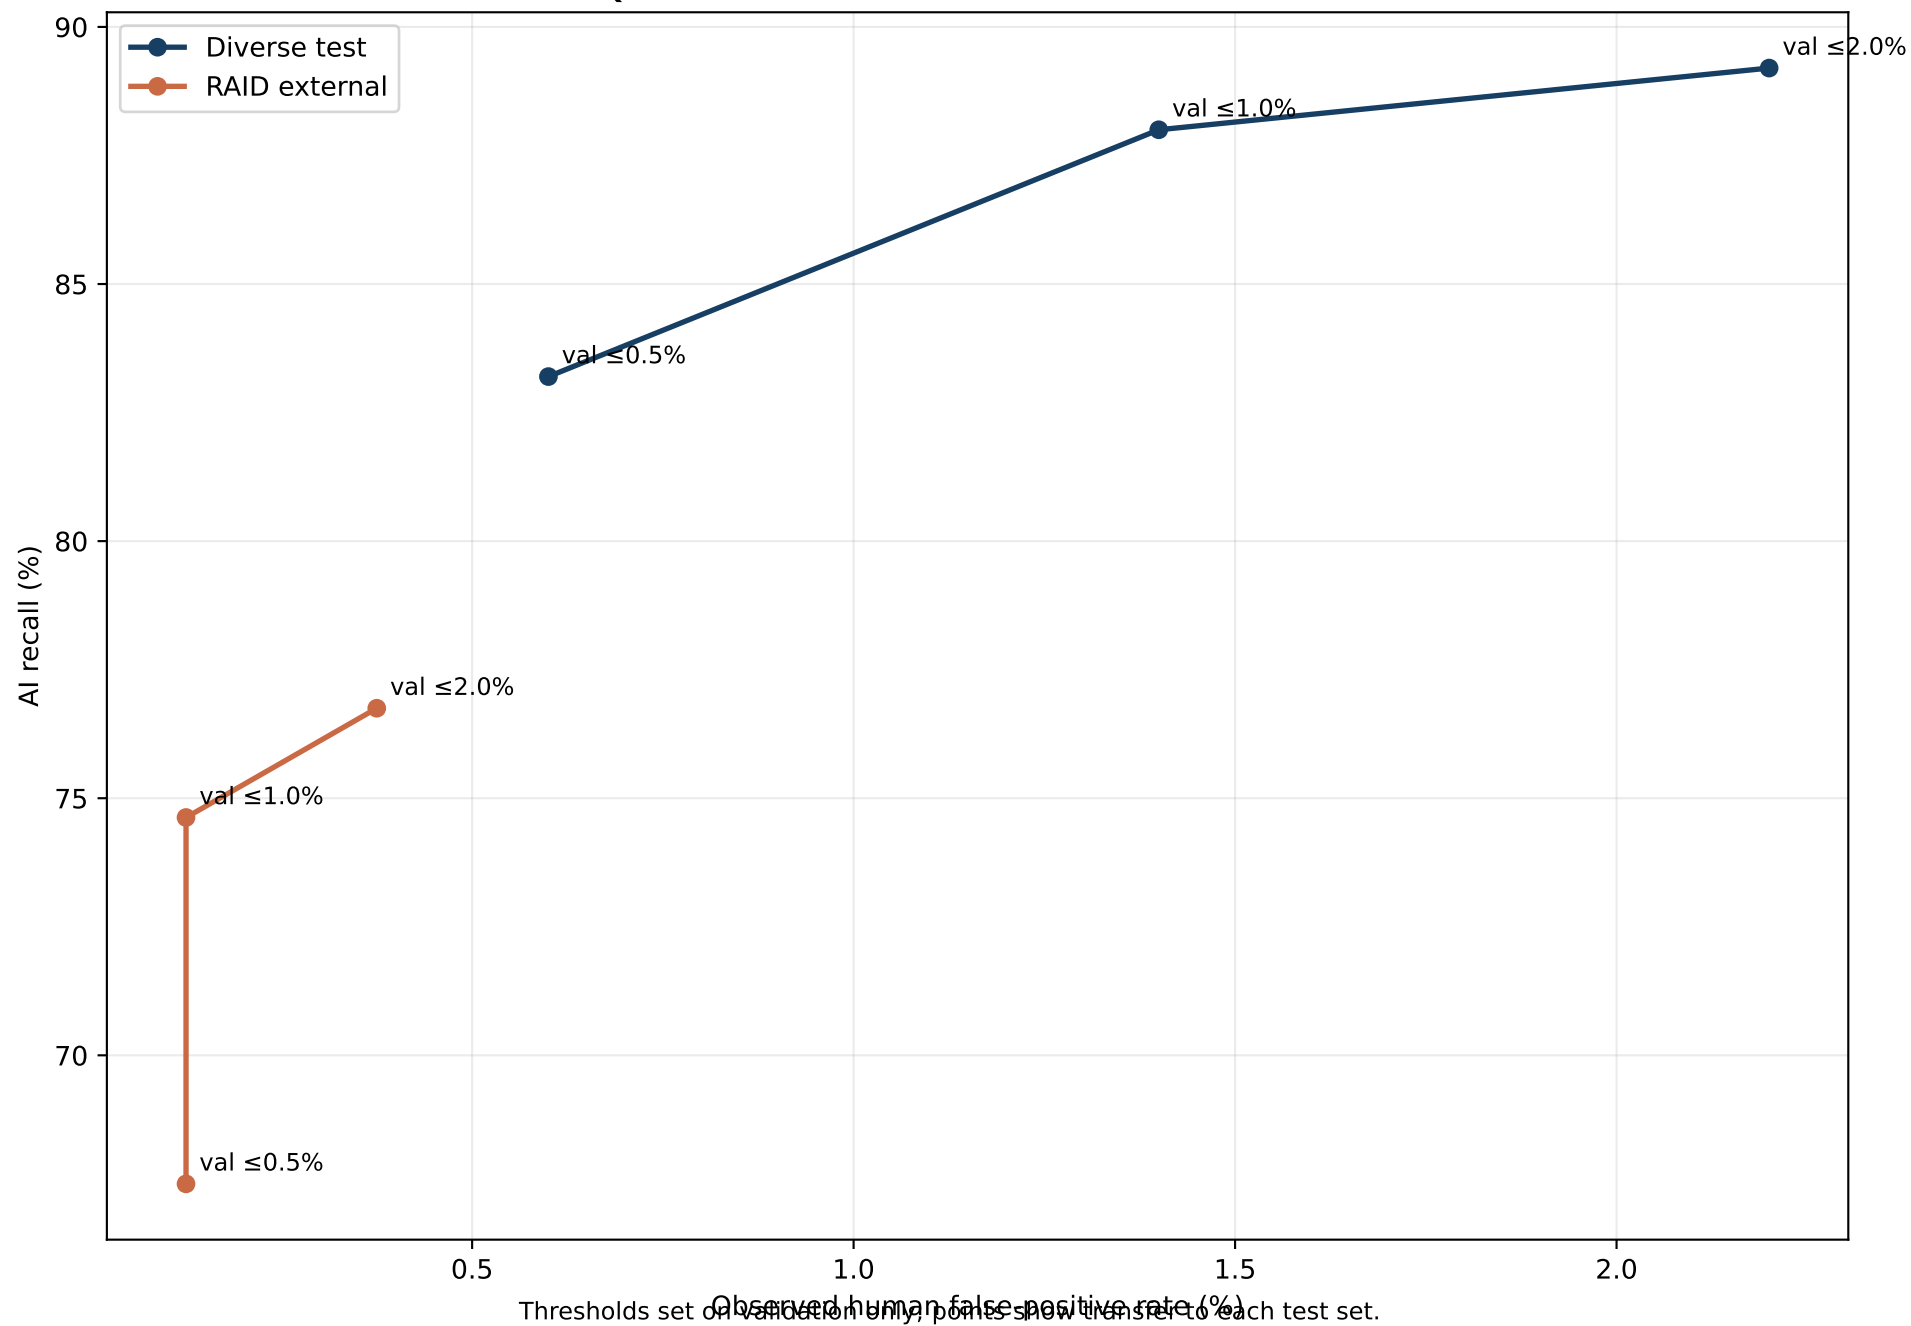


Model comparison across balanced tests

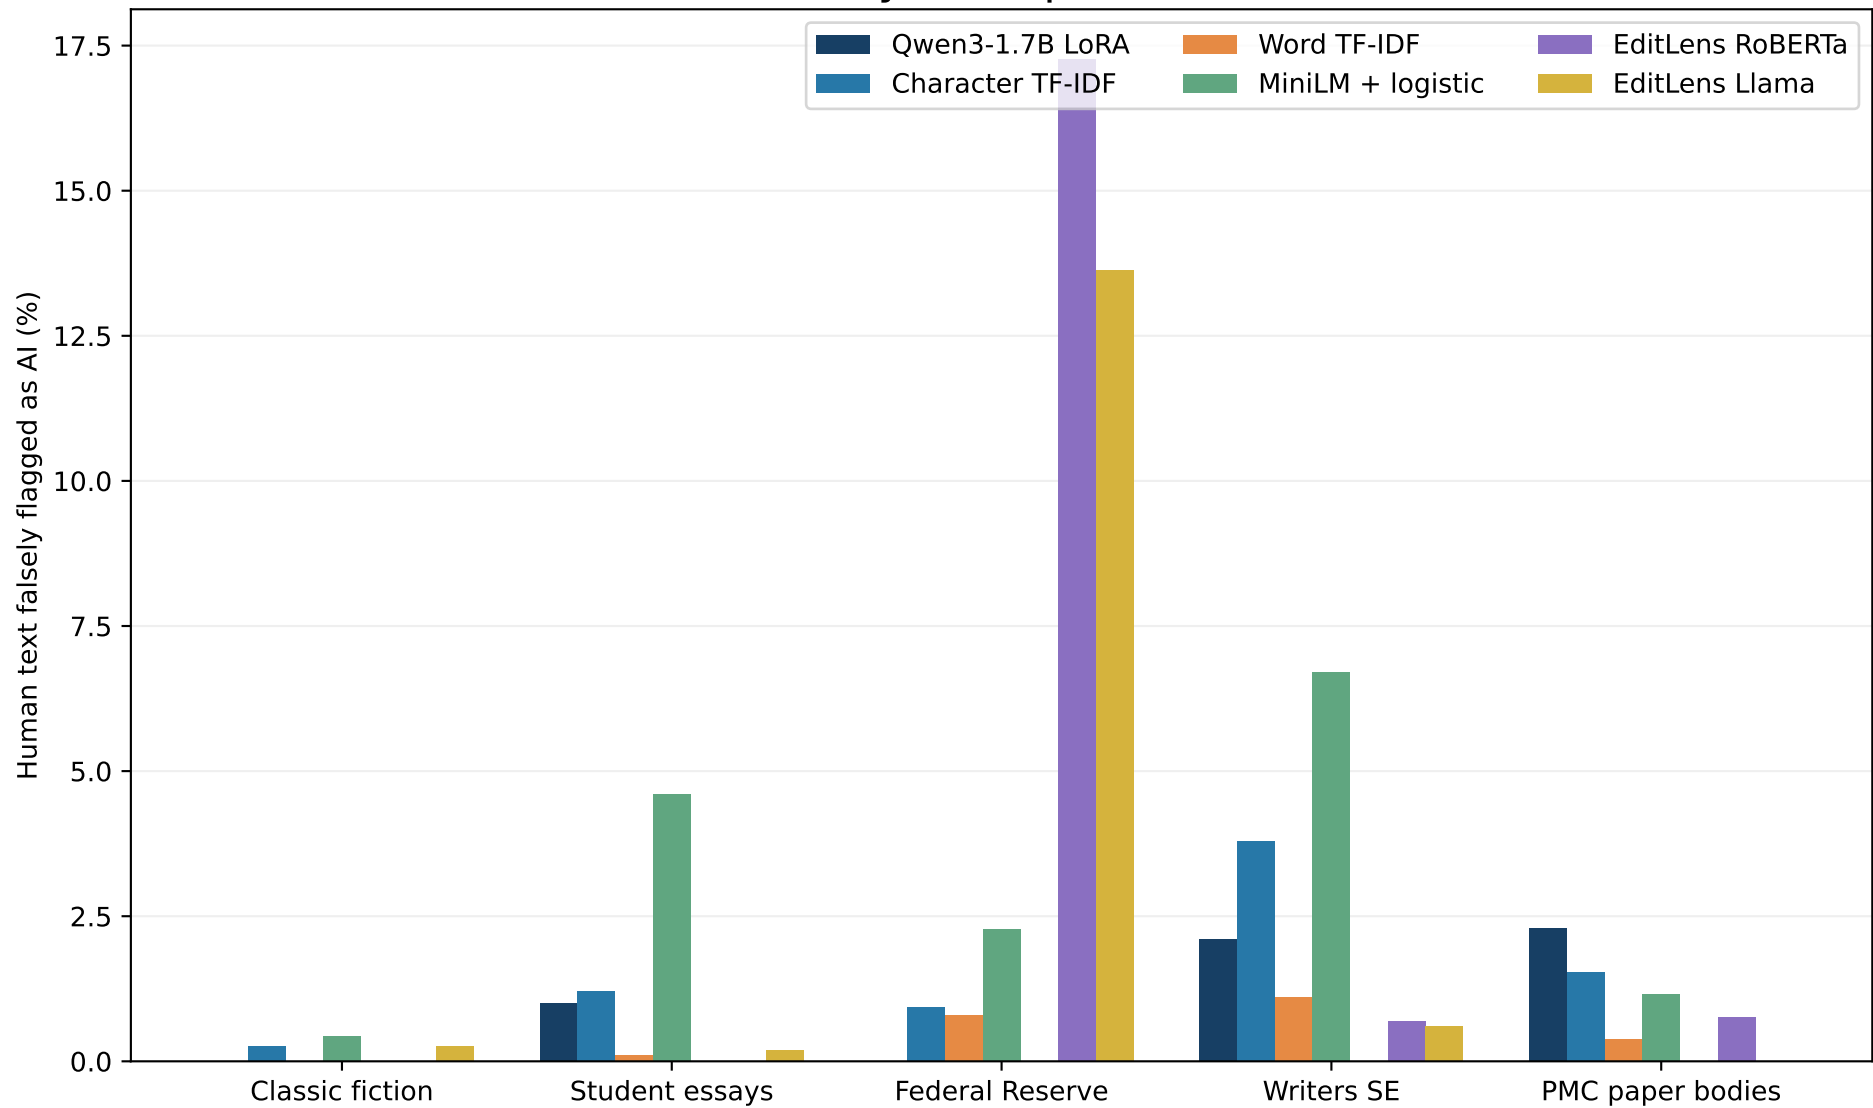
Model	Mixed AUC	Mixed FPR	Mixed recall	RAID AUC	RAID recall	Enron AUC	GPT-4 AUC	Paraphrase AUC
Qwen3-1.7B LoRA	0.989	2.2%	89.2%	0.974	76.8%	0.999	0.985	0.705
Character TF-IDF	0.882	1.8%	46.0%	0.826	25.5%	0.940	0.889	0.765
Word TF-IDF	0.862	1.0%	41.2%	0.803	24.6%	0.884	0.893	0.762
MiniLM + logistic	0.696	1.6%	9.6%	0.697	7.4%	0.813	0.742	0.691
Load Bearing cluster	0.481	3.6%	3.6%	0.513	4.6%	0.296	0.462	0.385
EditLens RoBERTa	0.814	3.6%	50.8%	0.842	37.8%	0.997	0.959	0.912
EditLens Llama	0.858	1.4%	59.2%	0.927	56.9%	1.000	0.991	0.968

Thresholds were selected from the diverse validation split at $\leq 2\%$ human FPR.
RAID is a separate source family; MAGE GPT-4/paraphrase shares a source family with training.
Load Bearing is a PR vocabulary-cluster model, included as a transfer baseline.

Qwen decision threshold trade-off



Human-only false-positive audits



Writers Stack Exchange is pre-2023 but individual human authorship is not verified.
PERSUADE is restricted student writing; all raw text remains outside Git.