



Our models and open Pangram at 2% human-calibration target

Shared human tests




External AI articles



Qwen detects an article when any token is highlighted; EditLens makes a document decision. Each model uses its own threshold from the same human calibration documents.


Additional human sources: our three checkpoints




ASAP test prompts are excluded from v10 training. Saved EditLens scores are unavailable for these sources.

Mixed-authorship localization at 2% generic calibration target

AI tokens highlighted




Human tokens falsely highlighted




Only documents containing both human and AI text. EditLens scores are broadcast from overlapping windows; Qwen predicts token labels directly.

External article ranking across thresholds

Full ROC




Low false-positive region



Decision thresholds selected on separate human calibration

Student essay transfer



Publication transfer

