



Our models and open Pangram at 2% human-calibration target

Shared human tests




LLMTrace pure AI



Qwen detects a document when any token is highlighted; EditLens makes a document decision. Each model uses its own threshold from the same human calibration documents.


Additional human sources: all five models




ASAP test prompts are excluded from v10 training. EditLens uses whole-document scores; Qwen flags any highlighted token.

Mixed-authorship localization at 2% generic calibration target

AI tokens highlighted




Human tokens falsely highlighted




Only documents containing both human and AI text. EditLens scores are broadcast from overlapping windows; Qwen predicts token labels directly.

External article ranking across thresholds



Full ROC



Low false-positive region




LLMTrace pure-document ranking across thresholds



LLMTrace contributes some v8-v10 training records; evaluation IDs are disjoint but this is a same-source test.

Decision thresholds selected on separate human calibration

Student essay transfer



Publication transfer

