



Pure documents: source shift changes the ranking

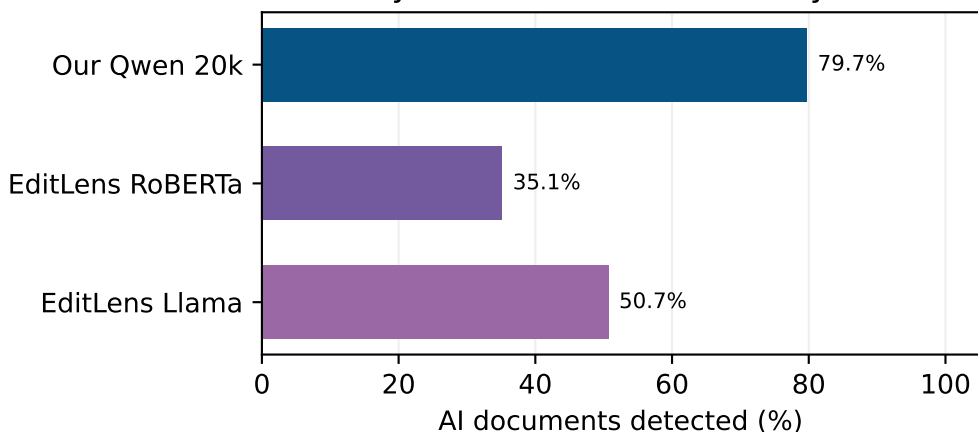
LLMTrace test: fully AI



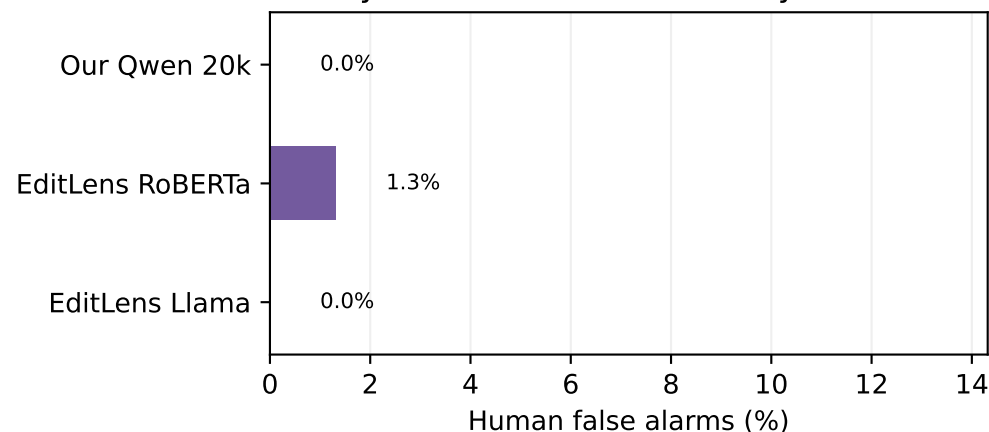
LLMTrace test: fully human



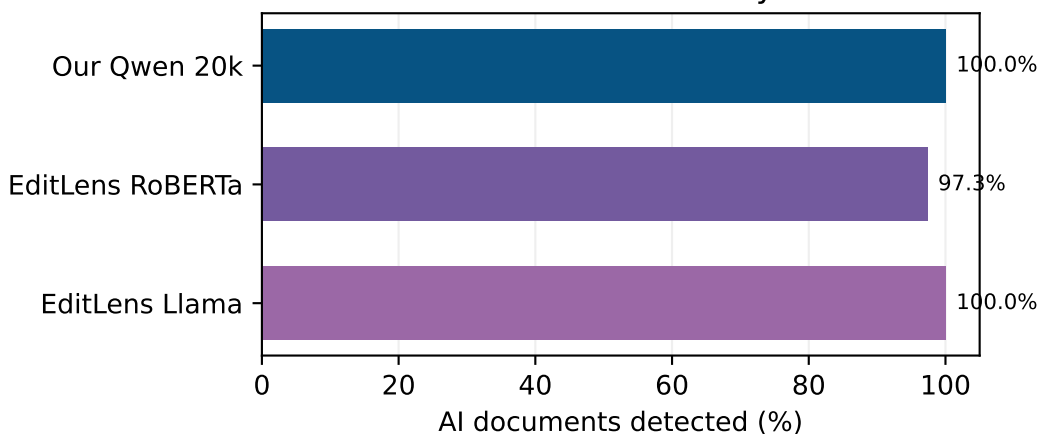
Synthetic v4 validation: fully AI



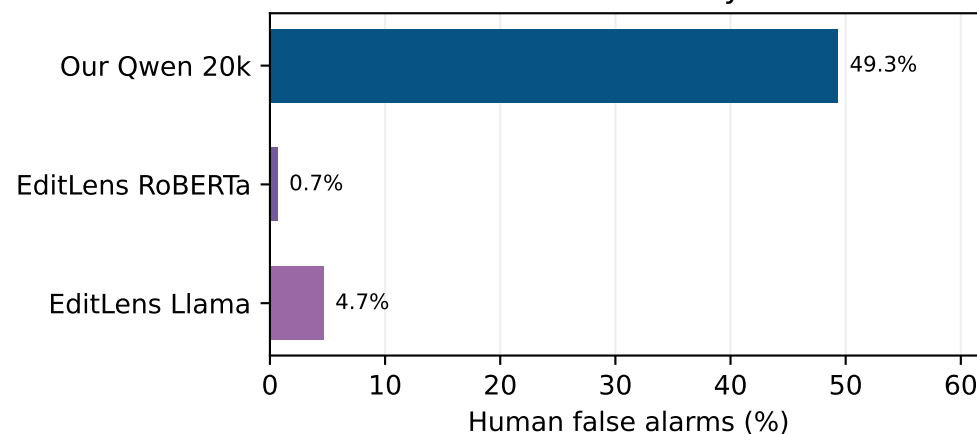
Synthetic v4 validation: fully human



External articles: fully AI



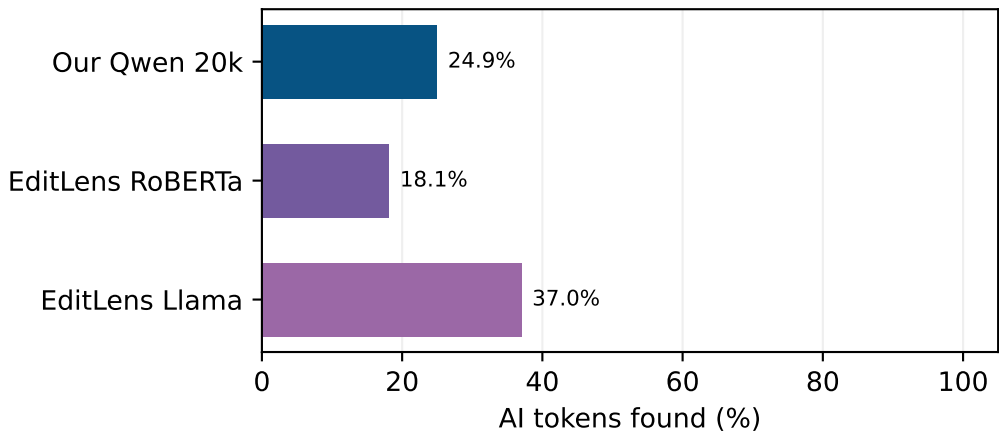
External articles: fully human



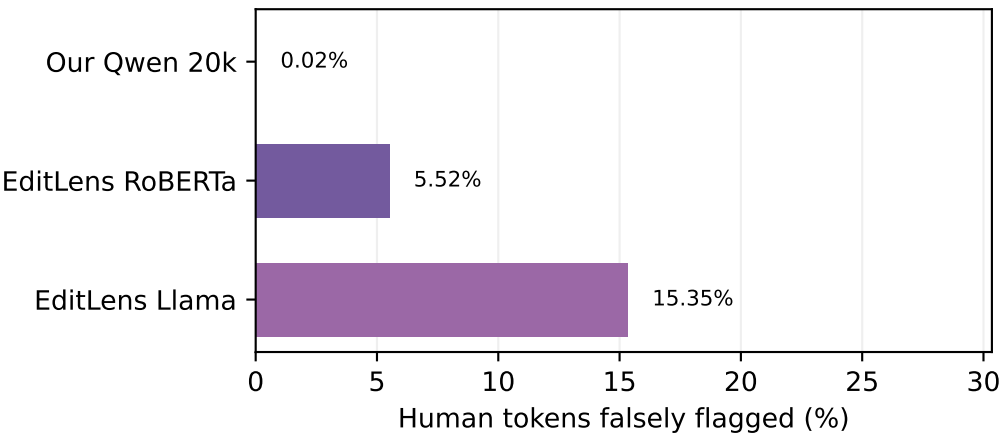
Thresholds fixed on separate human calibration. LLMTrace is an in-corpus holdout for our model; external articles are a different source.

Mixed documents: token-level performance

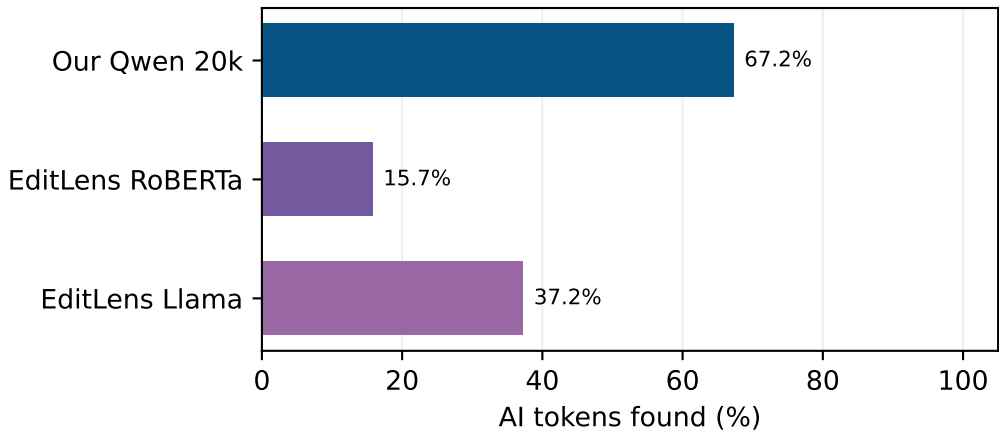
LLMTrace test: AI-token recall



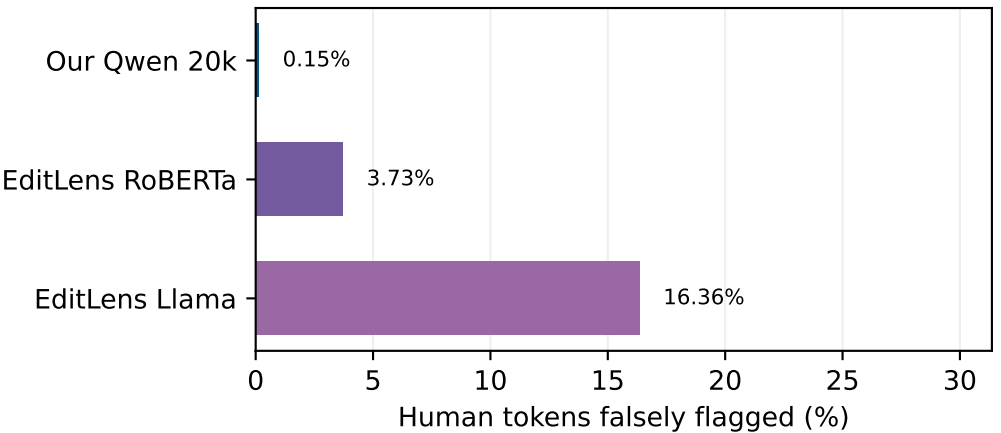
LLMTrace test: human-token FPR



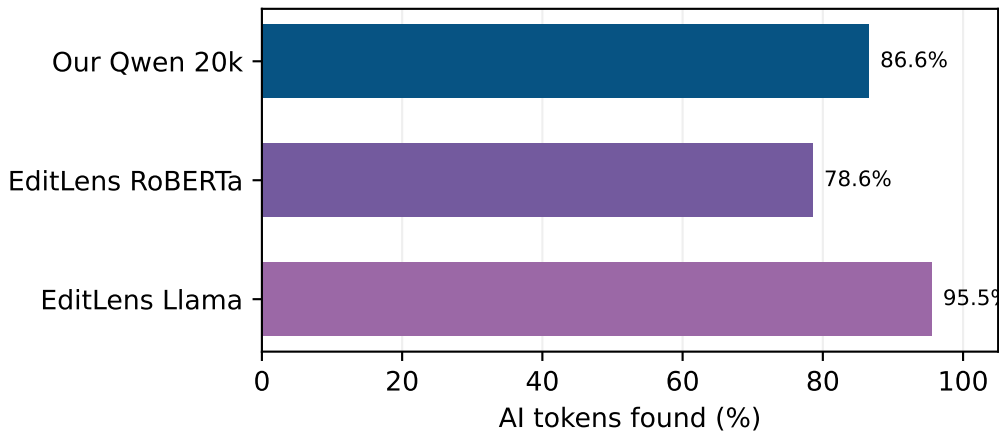
Synthetic v4 validation: AI-token recall



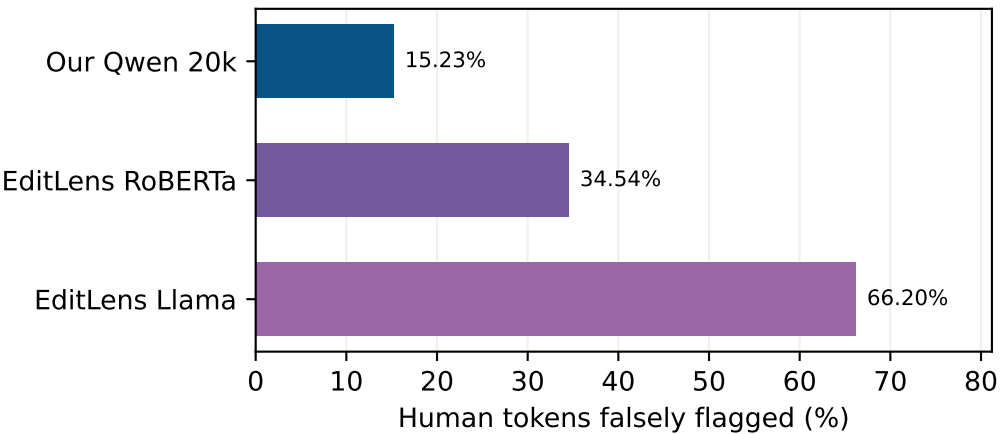
Synthetic v4 validation: human-token FPR



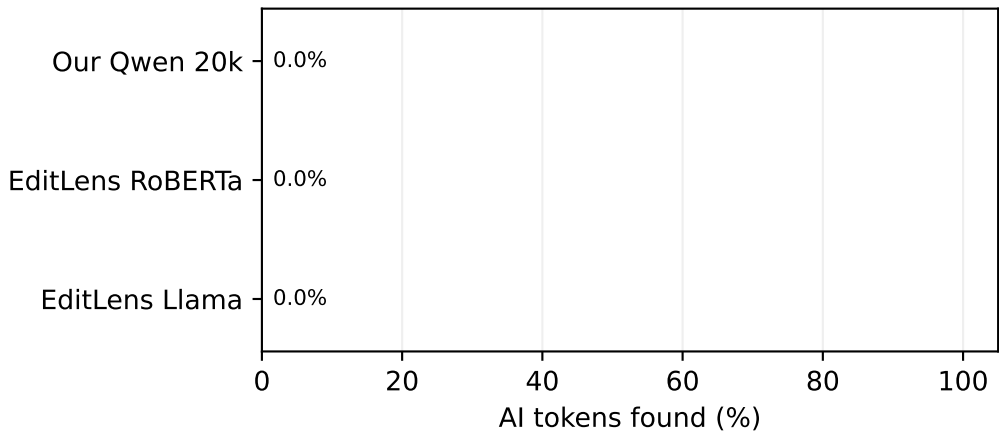
AITDNA collaboration: AI-token recall



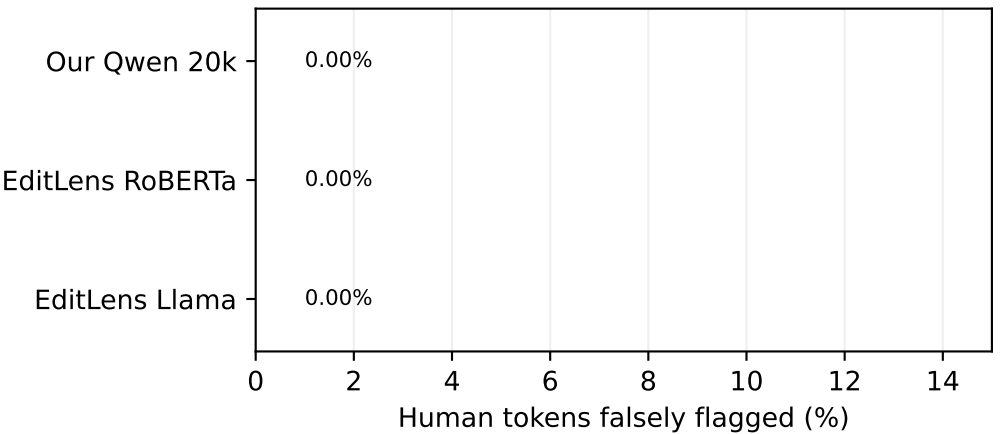
AITDNA collaboration: human-token FPR



CoAuthor collaboration: AI-token recall

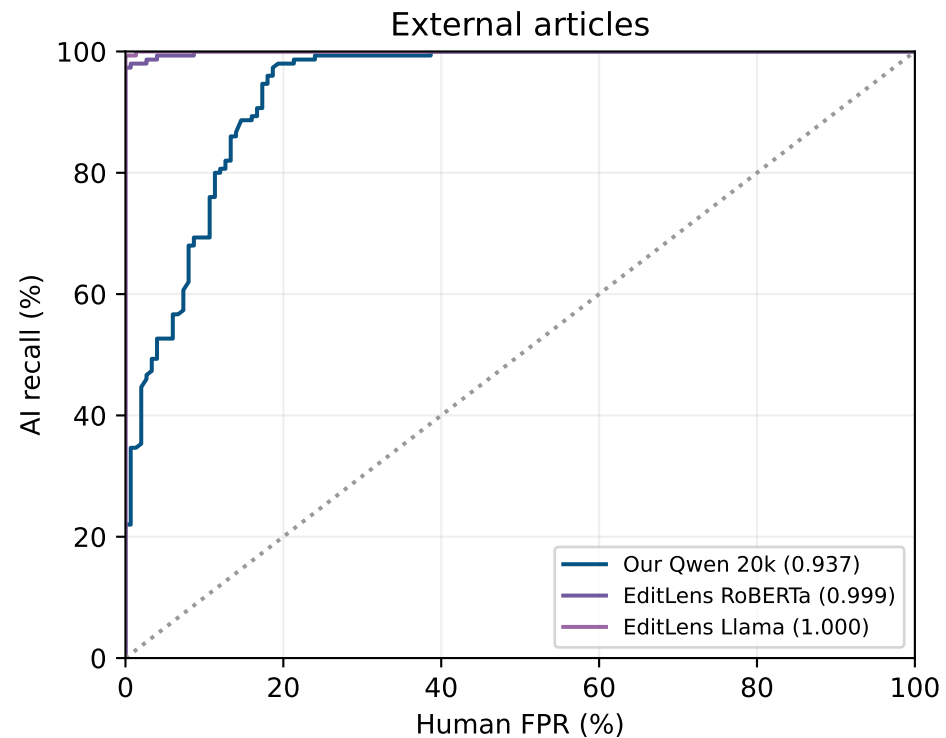
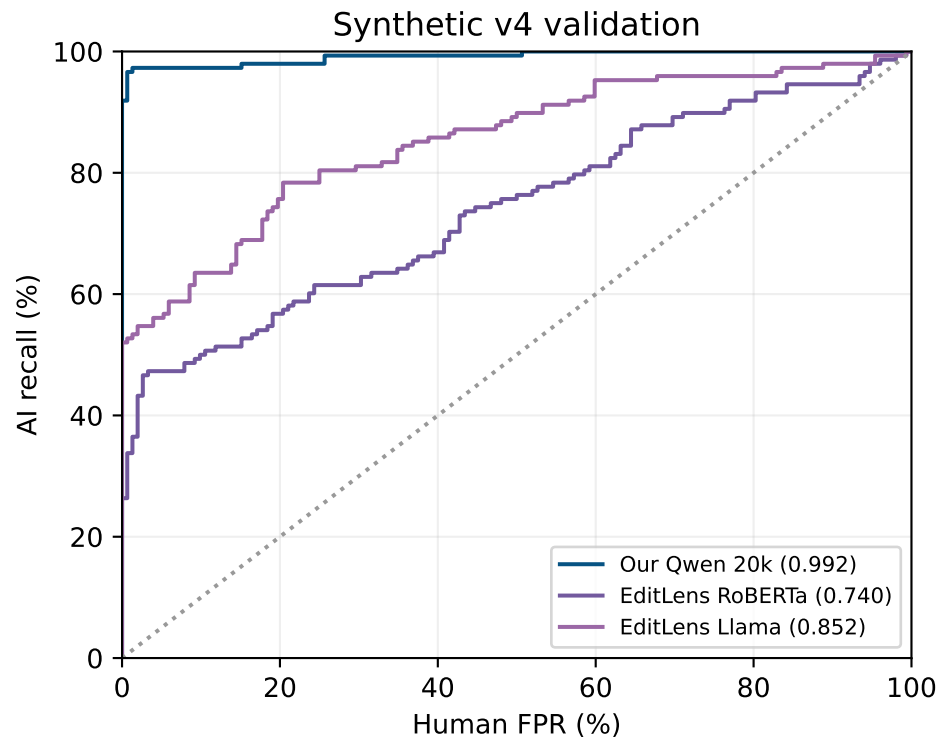
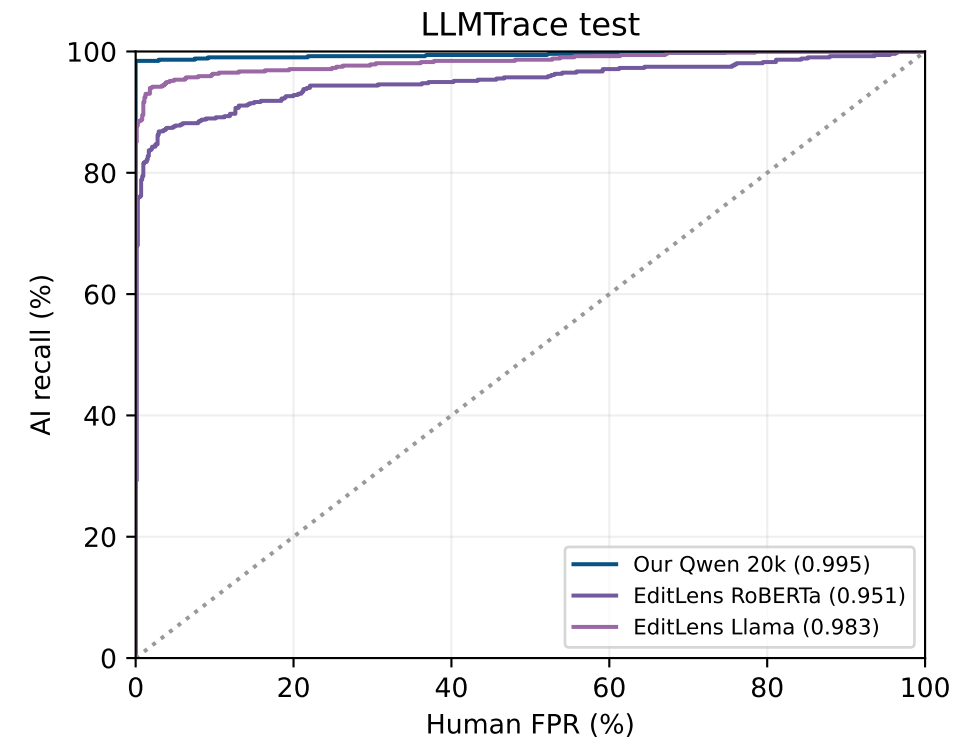


CoAuthor collaboration: human-token FPR



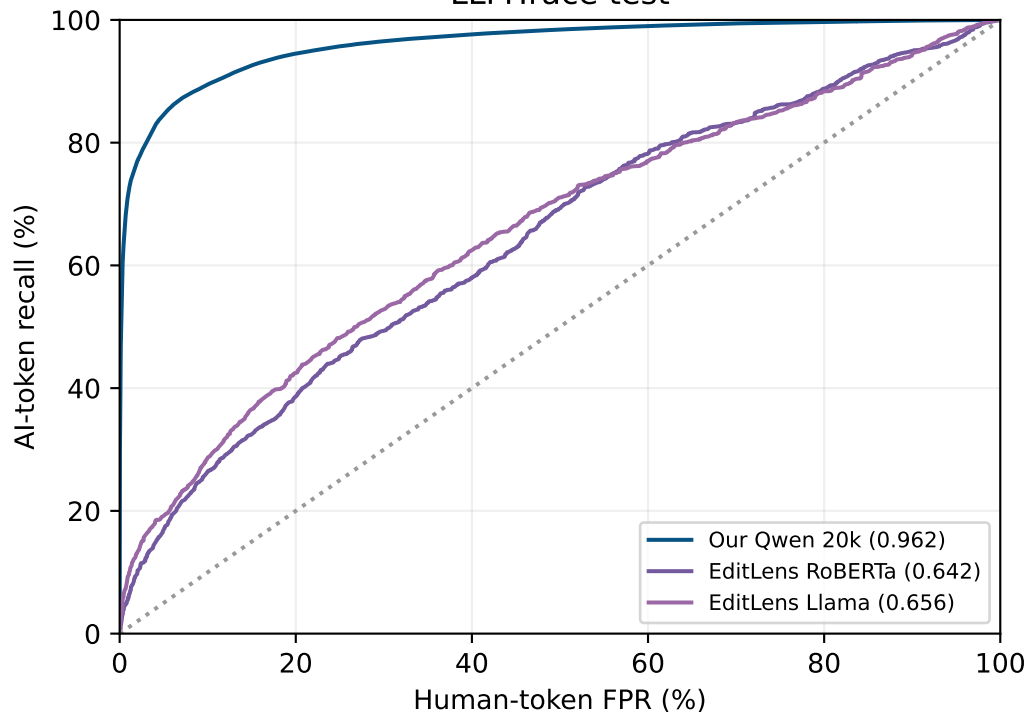
EditLens predictions are broadcast window scores; their spans are coarse. LLMTrace is an in-corpus holdout for our model.

Document ROC curves (AUROC in legend)




Mixed-document token ROC curves (AUROC in legend)


LLMTrace test



Synthetic v4 validation



AITDNA collaboration



CoAuthor collaboration

