




Hyperparameter pilot: small gains, useful candidates

24 trials · 7 completed full budget · 17 pruned · 0 Ray failures

Same 800-row validation set used throughout; results are not held-out test estimates.

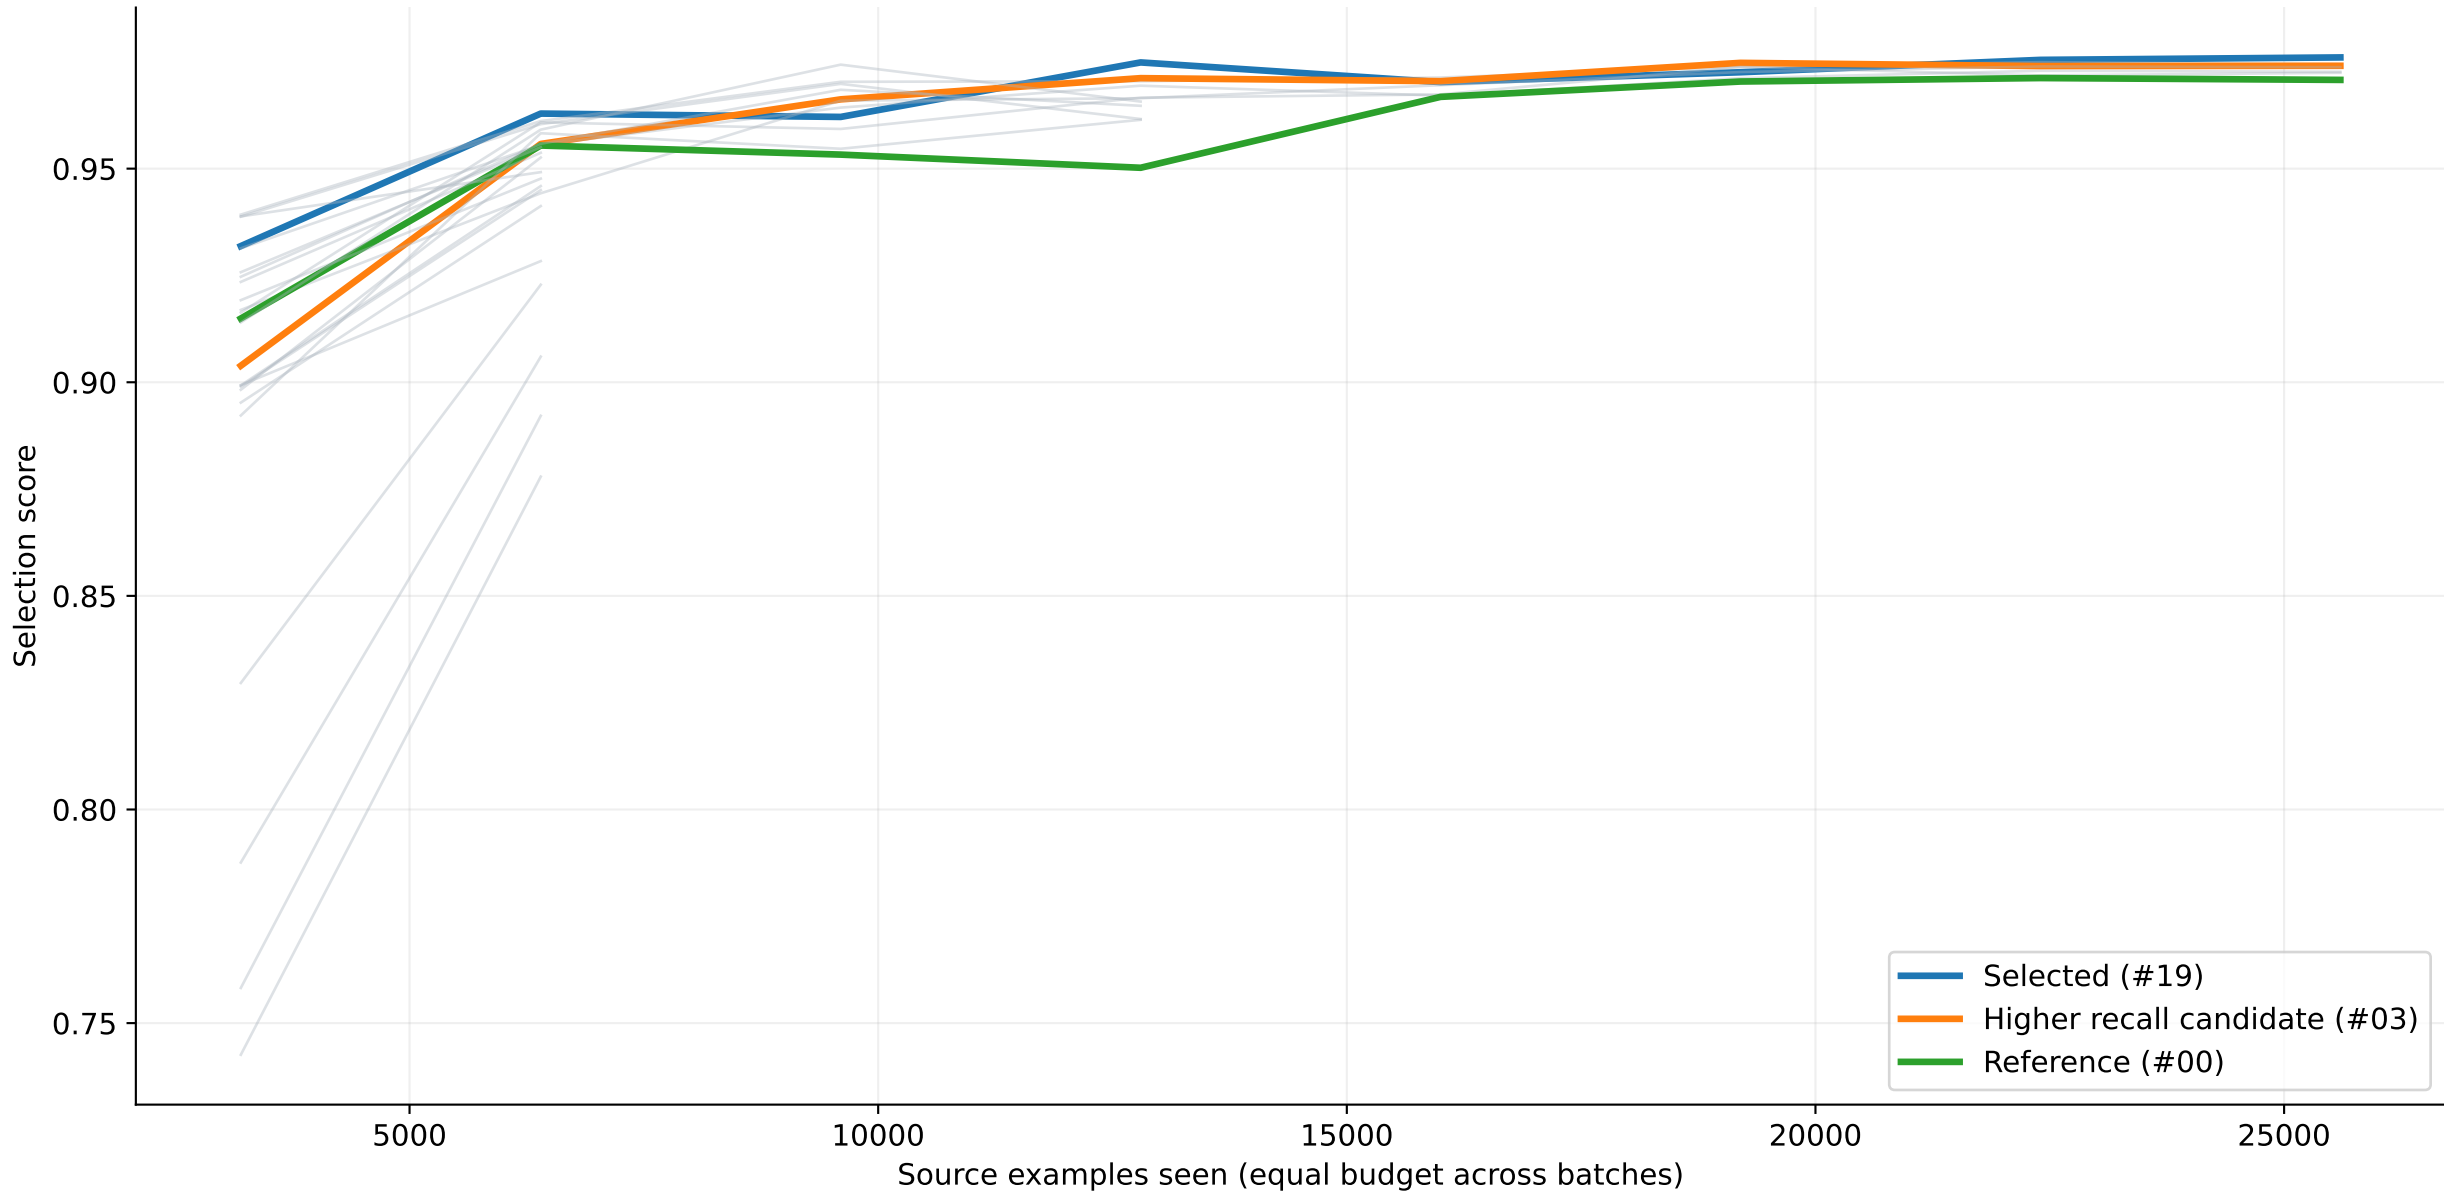


Selected: LR $7.61e-5$ · effective batch 8 · rank 32 / alpha 64 · dropout 0.0688

Recall gain: 3 of 400 AI examples. Worst-category recall: 85.0%, versus reference 87.5%.

Selection score = 70% standardized partial AUROC + 30% full AUROC.

Validation score as training progresses



Lines ending early were pruned by ASHA; do not interpret them as full-budget results.

All trials: full-budget candidates first

Trial	Status	Examples	LR	Batch	Rank	Dropout	pAUC	Recall %
00019	completed	25600	7.61e-05	8	32	0.069	0.9681	93.50
00007	completed	25600	1.01e-04	8	16	0.114	0.9665	93.50
00003	completed	25600	8.65e-05	32	32	0.019	0.9665	95.00
00001	completed	25600	8.91e-05	16	16	0.129	0.9658	93.50
00015	completed	25600	7.52e-05	8	8	0.125	0.9645	92.50
00005	completed	25600	6.61e-05	16	32	0.067	0.9637	92.75
00000	completed	25600	5.00e-05	16	16	0.100	0.9628	92.75
00008	pruned	12800	3.39e-05	8	32	0.134	0.9567	92.00
00012	pruned	12800	3.52e-05	8	16	0.028	0.9560	92.50
00011	pruned	12800	8.33e-05	16	32	0.049	0.9501	91.25
00016	pruned	12800	9.57e-05	32	16	0.043	0.9493	90.25
00002	pruned	6400	7.47e-05	8	8	0.146	0.9433	90.75
00021	pruned	6400	6.99e-05	16	16	0.085	0.9406	90.00
00017	pruned	6400	7.22e-05	16	8	0.030	0.9397	88.75
00022	pruned	6400	8.73e-05	16	16	0.083	0.9344	89.00
00020	pruned	6400	5.56e-05	8	8	0.017	0.9327	88.00
00014	pruned	6400	7.01e-05	32	16	0.125	0.9305	85.75
00004	pruned	6400	4.23e-05	16	16	0.139	0.9321	86.50
00009	pruned	6400	9.00e-05	32	8	0.070	0.9252	85.00
00006	pruned	6400	2.53e-05	8	16	0.010	0.9076	84.50
00023	pruned	6400	5.44e-05	32	8	0.005	0.9037	80.50
00018	pruned	6400	1.53e-05	32	32	0.100	0.8840	80.75
00013	pruned	6400	2.02e-05	32	16	0.034	0.8679	75.25
00010	pruned	6400	1.66e-05	16	8	0.102	0.8527	73.00

Pruned results use less training and are not directly comparable to completed trials.
Recall uses a threshold chosen on these same validation humans; it is not a guaranteed deployment FPR.