



Pure-document detection: our size sweep and Open Pangram




Held-out LLMTrace: 516 AI + 720 human documents. Our 20k run trained on 15,036 other LLMTrace records.
Thresholds use the same separate 1,120 human docs ($\leq 5\%$ FPR). Ours needs any AI token; EditLens uses mean window score.

Threshold-independent ranking and held-out human false alarms

Ranking quality on the pure held-out set



Additional human-only stress test



AUROC uses each model's document score: maximum token score for ours, mean window score for EditLens.
Zero observed false alarms is an estimate on these datasets, not a guarantee for arbitrary human text.