



Pure-document detection: our size sweep and Open Pangram

AI-document recall



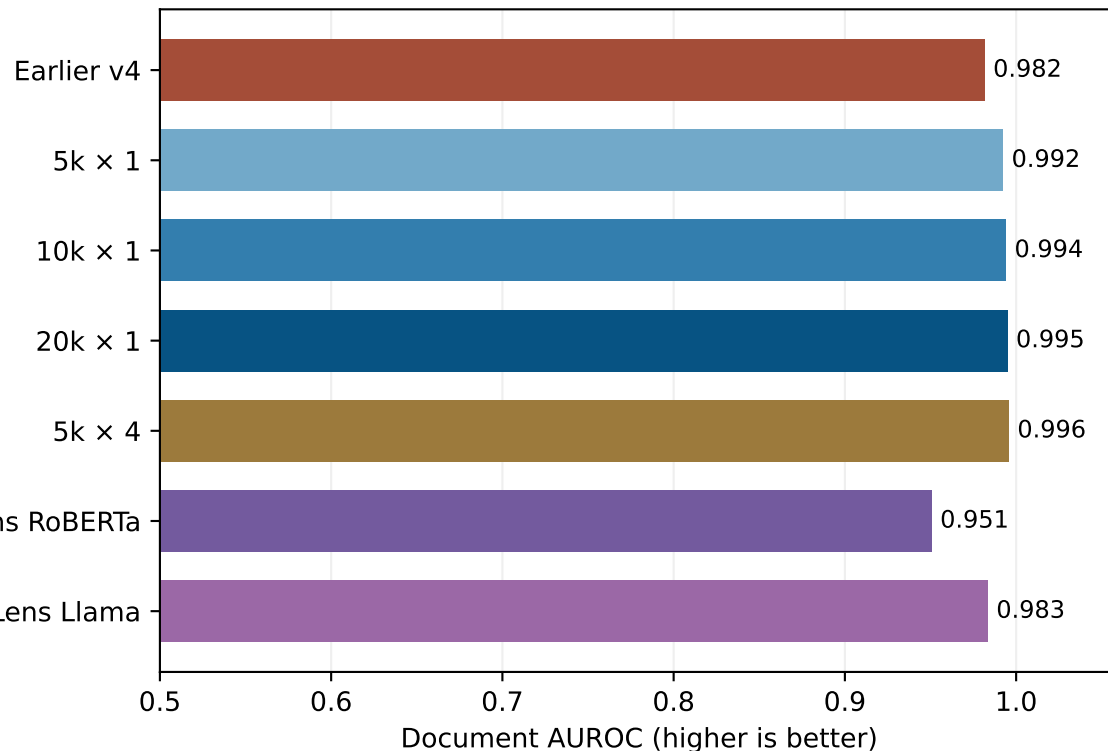
Human-document false-positive rate



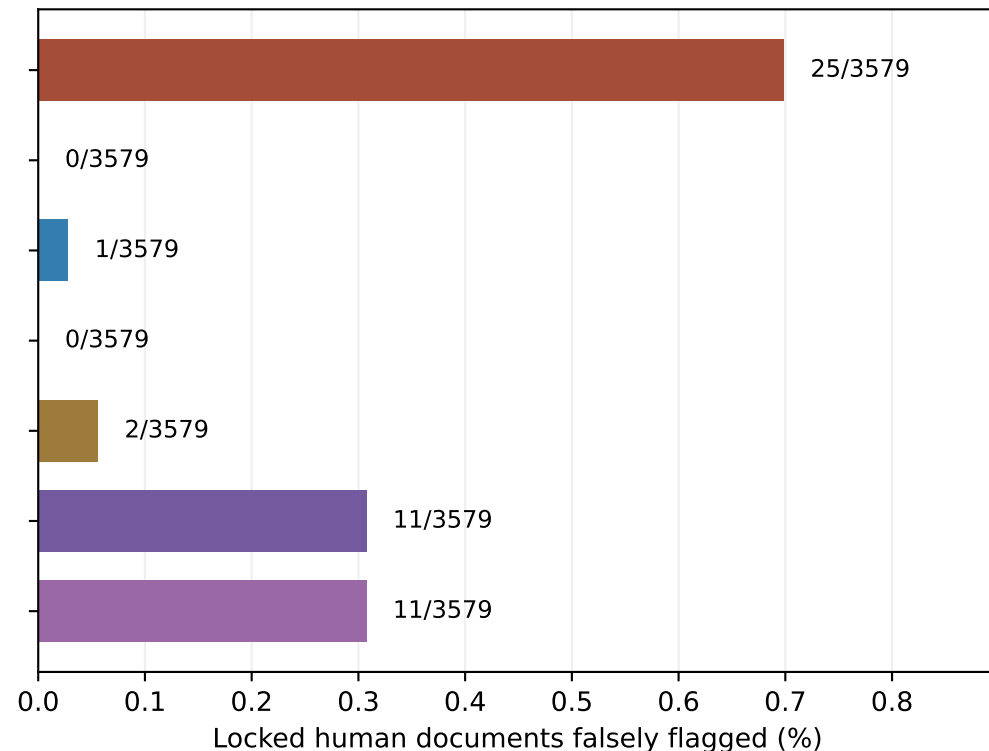
Held-out LLMTrace: 516 AI + 720 human documents. All thresholds set on the same separate 1,120 human documents with $\leq 5\%$ false-alarm allowance. Our detector flags a document if any token is highlighted; EditLens averages native-window scores.

Threshold-independent ranking and held-out human false alarms

Ranking quality on the pure held-out set



Additional human-only stress test



AUROC uses each model's document score: maximum token score for ours, mean window score for EditLens.
Zero observed false alarms is an estimate on these datasets, not a guarantee for arbitrary human text.