


1. Whole-document decisions


At the frozen threshold. A false alarm means a fully human document was flagged as AI.



Verdict: our Qwen catches article AI, but flags 74 of 150 human articles. That false-alarm rate is unacceptable.

2. Can the models highlight AI passages in mixed writing?



Read each row as “out of 100”. AITDNA is real collaboration; LLMTrace/synthetic are constructed; CoAuthor has short insertions.



Verdict: Qwen is useful on AITDNA, misses most LLMTrace AI text, and misses all CoAuthor AI at this threshold.

3. Threshold-free document ROC curves


A curve near the upper-left is better. The dots mark the frozen operating threshold used on pages 1-2.



The article ROC also shows ranking weakness: AUC 0.937 for Qwen versus 0.999/1.000 for the baselines. Moving the threshold along will trade away AI recall; source-balanced retraining is the next test.

4. Human-publication false alarms

Our Qwen flags attributed-human published articles across several publications, not just one source.



These articles have named human authors; their production workflows were not independently verified as AI-free. This 150-article set stays frozen for evaluation. It will not be used to set a new threshold or as training text.