



Held-out LLMTrace: separate the fully human and fully AI documents

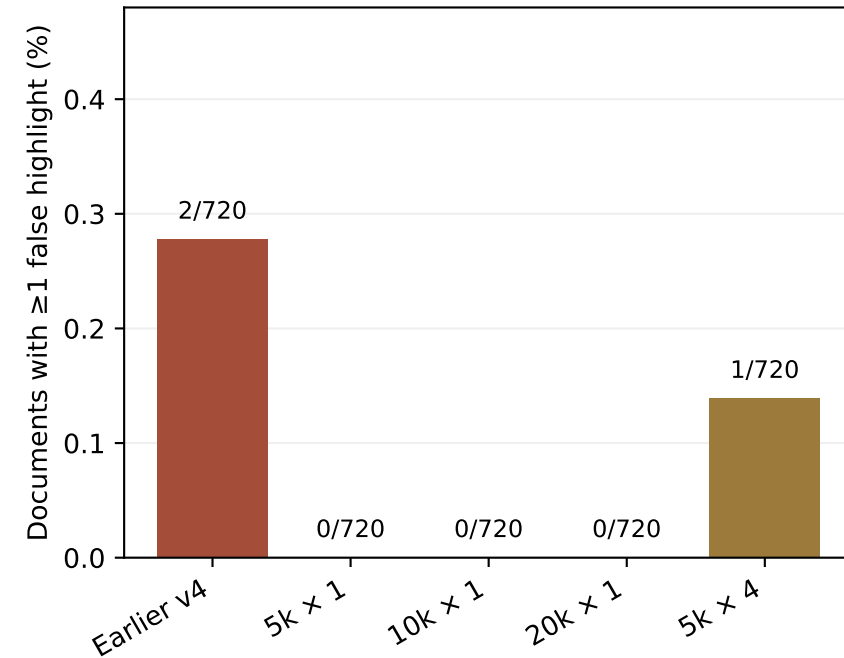
Fully AI: tokens highlighted



Fully AI: documents detected

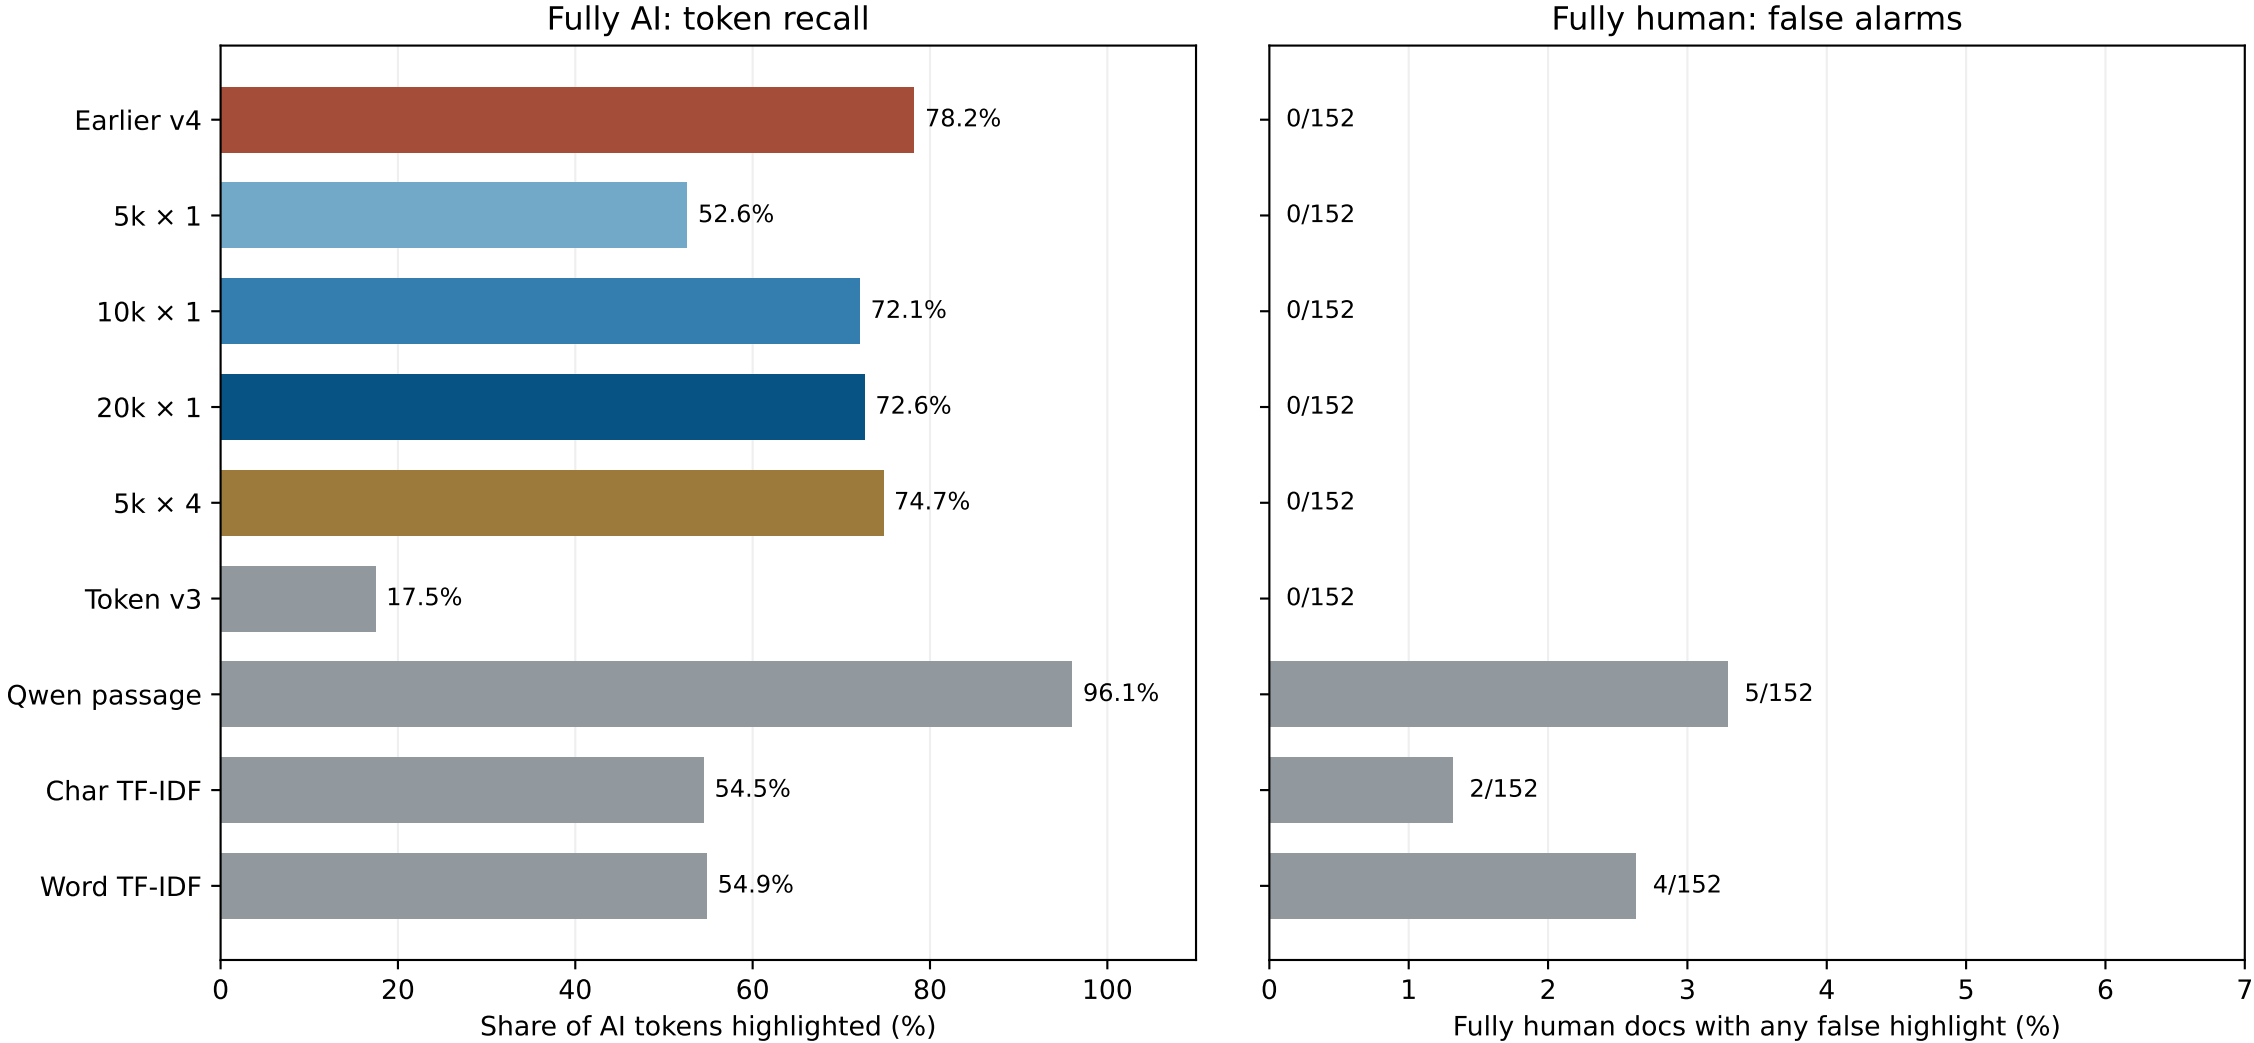


Fully human: false alarms



516 fully AI and 720 fully human documents; nine domains. Each threshold was set on a separate 1,120-document human calibration set. Detecting ≥ 1 token is a permissive document rule; token recall measures how much AI text was actually highlighted.

Shared v4 synthetic evaluation: size sweep versus baselines



Same 148 fully AI and 152 fully human documents for every model. Each model uses its own threshold calibrated to 5% document-any false highlight on the same separate human calibration set.