


Overall comparison

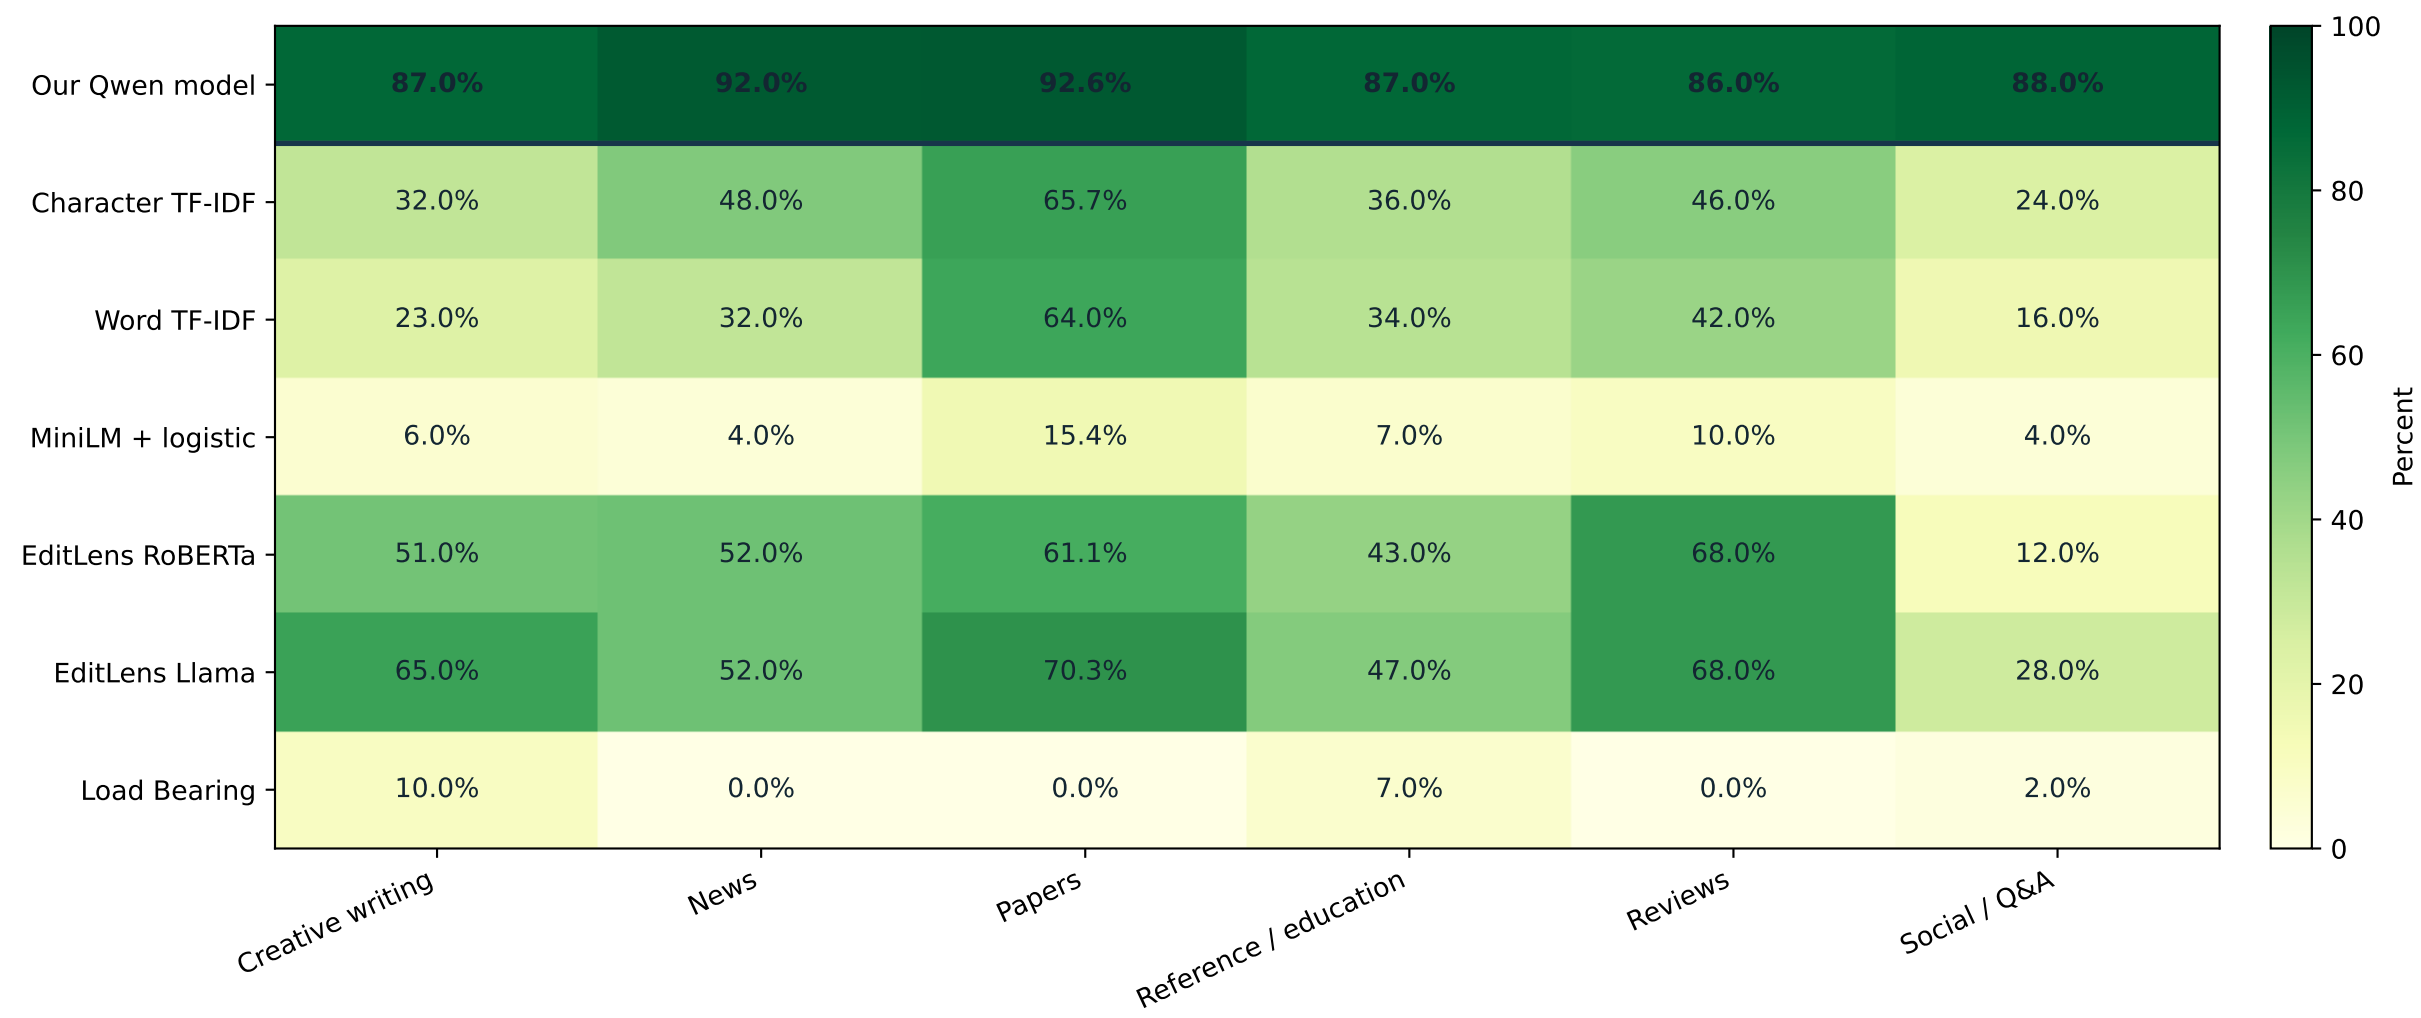
Equal-weight mean of the 1,000-row diverse test and 1,600-row RAID test. Both are 50% human / 50% AI.



Recall: share of AI passages caught. Balanced accuracy: mean of AI recall and human specificity. False positives: share of human passages wrongly flagged.
Each model uses its own threshold chosen on the same validation set for $\leq 2\%$ human false positives. Observed test rates can differ.

Diverse test: AI recall by category

Higher is better. Each cell is the percent of AI passages correctly flagged.

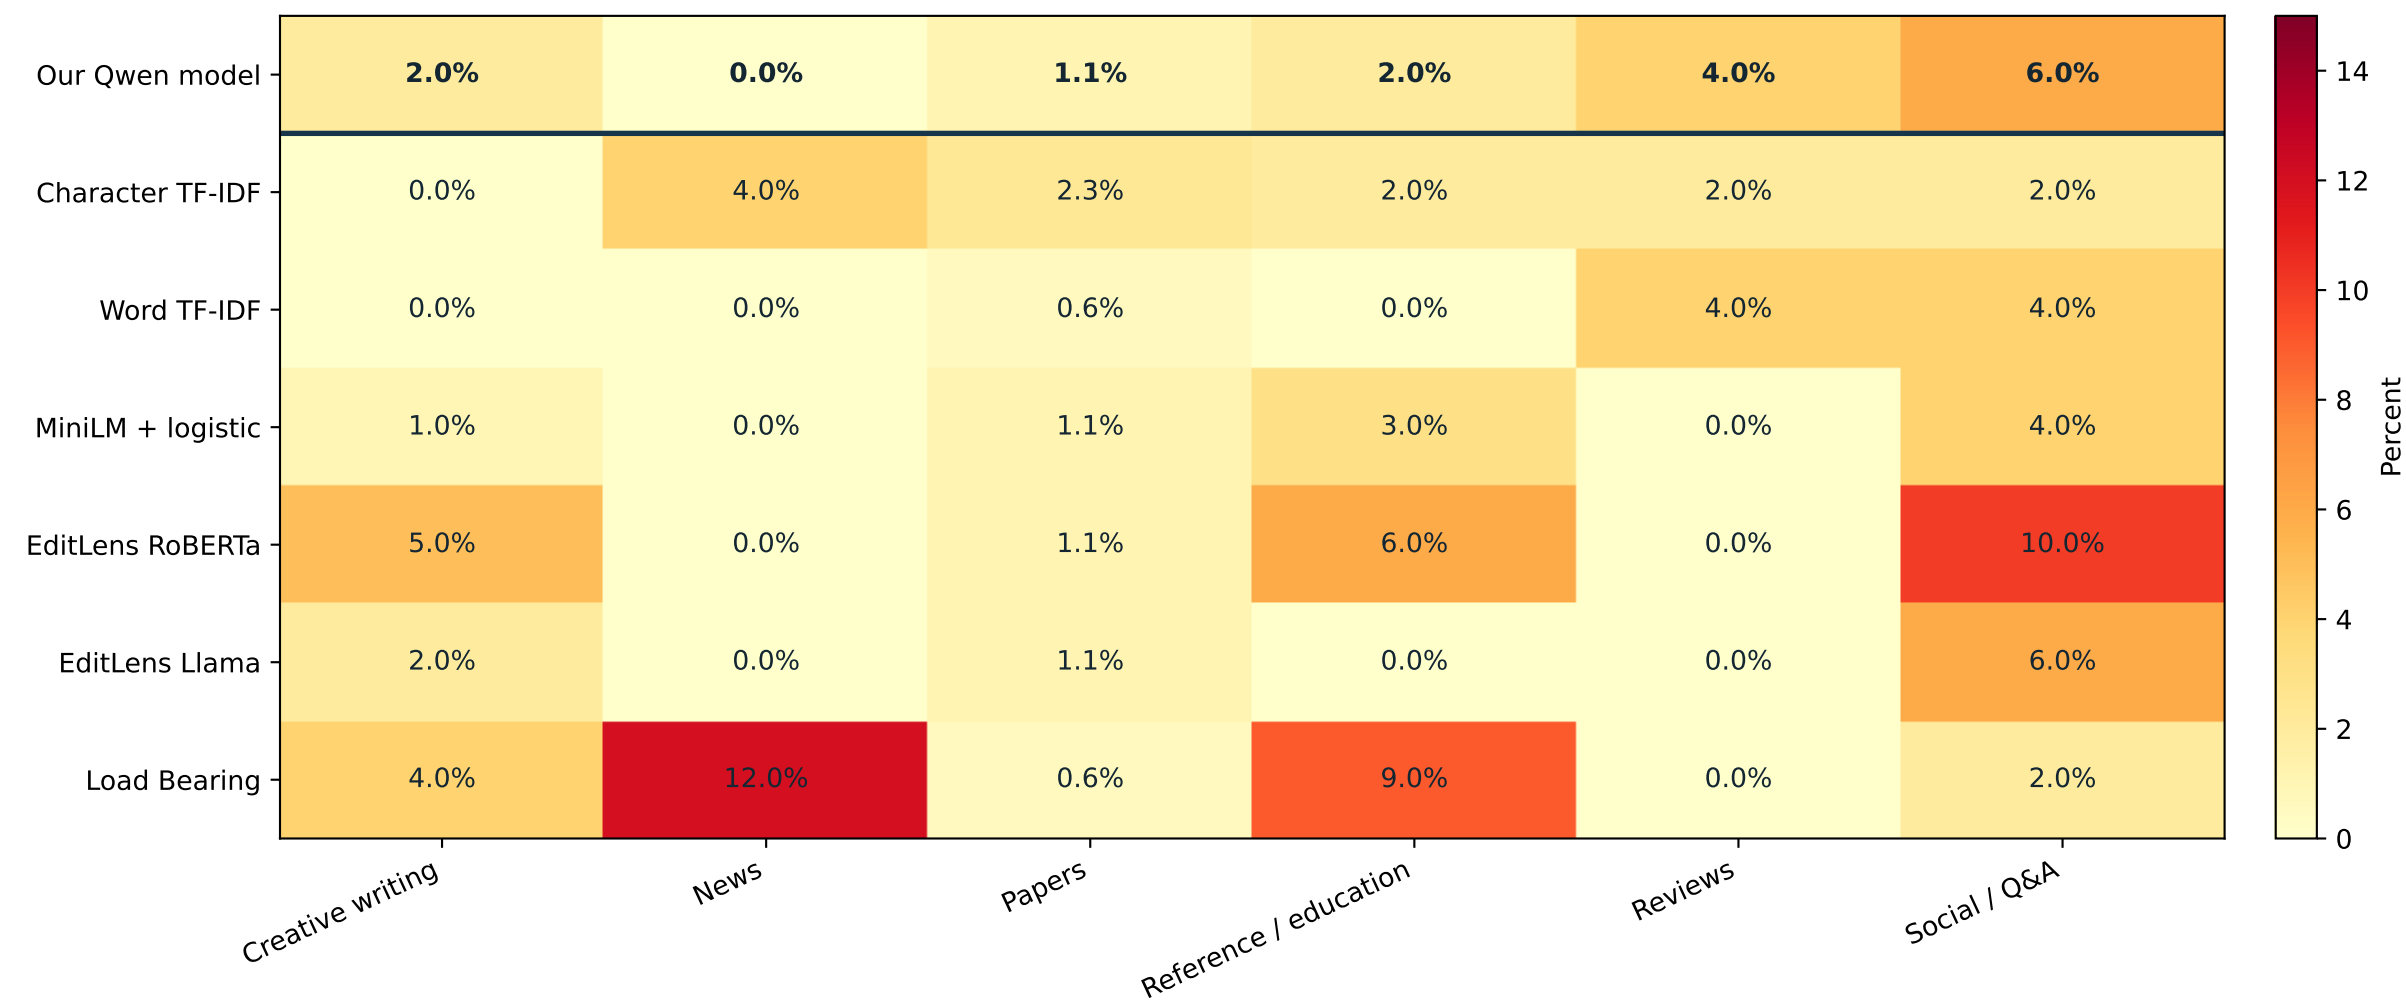


Passages per category (AI): Creative writing 100, News 25, Papers 175, Reference / education 100, Reviews 50, Social / Q&A 50

Small categories have noisy percentages; one error changes a 50-passage category by 2 points.

Diverse test: human false positives by category

Lower is better. Each cell is the percent of human passages wrongly flagged.

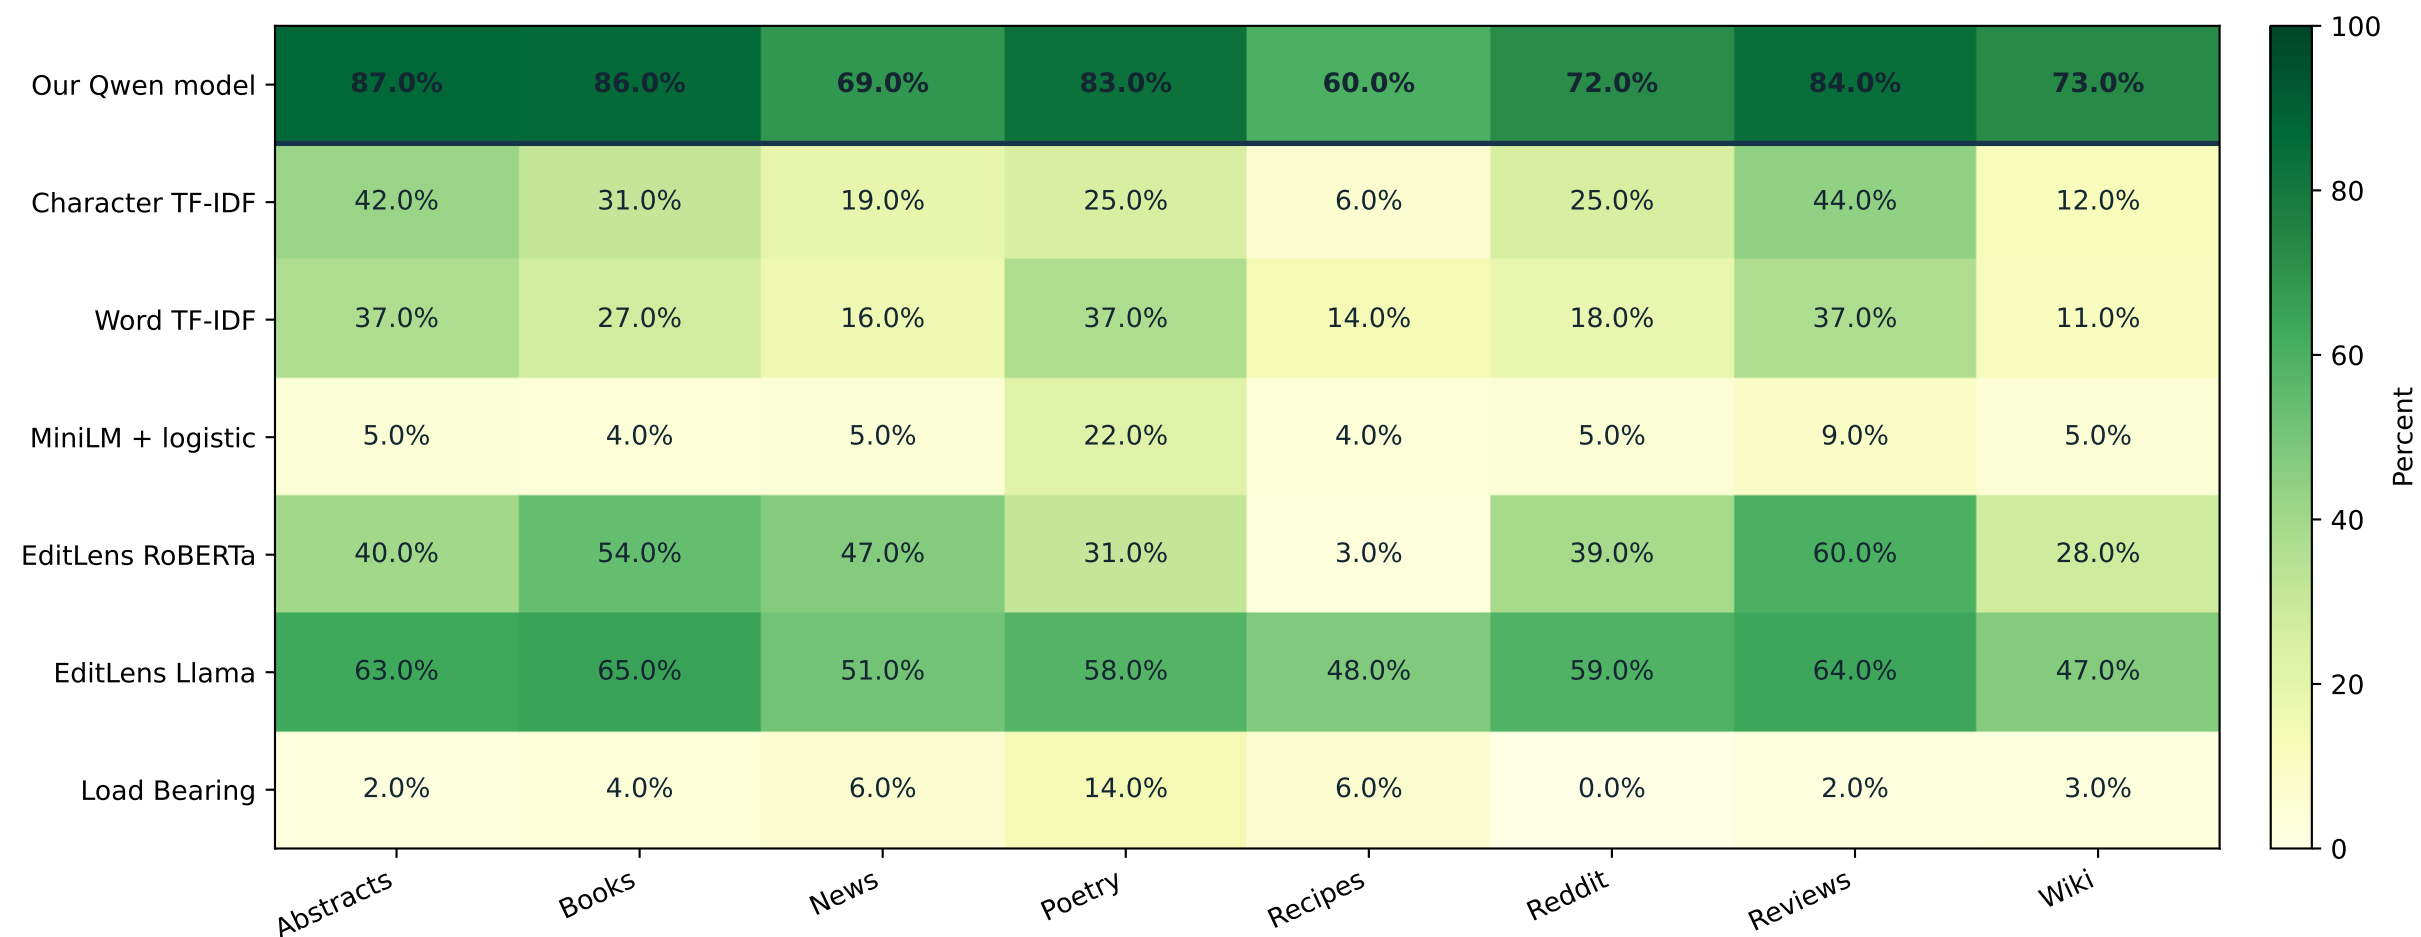


Passages per category (human): Creative writing 100, News 25, Papers 175, Reference / education 100, Reviews 50, Social / Q&A 50

Small categories have noisy percentages; one error changes a 50-passage category by 2 points.

RAID independent test: AI recall by category

Higher is better. Each cell is the percent of AI passages correctly flagged.

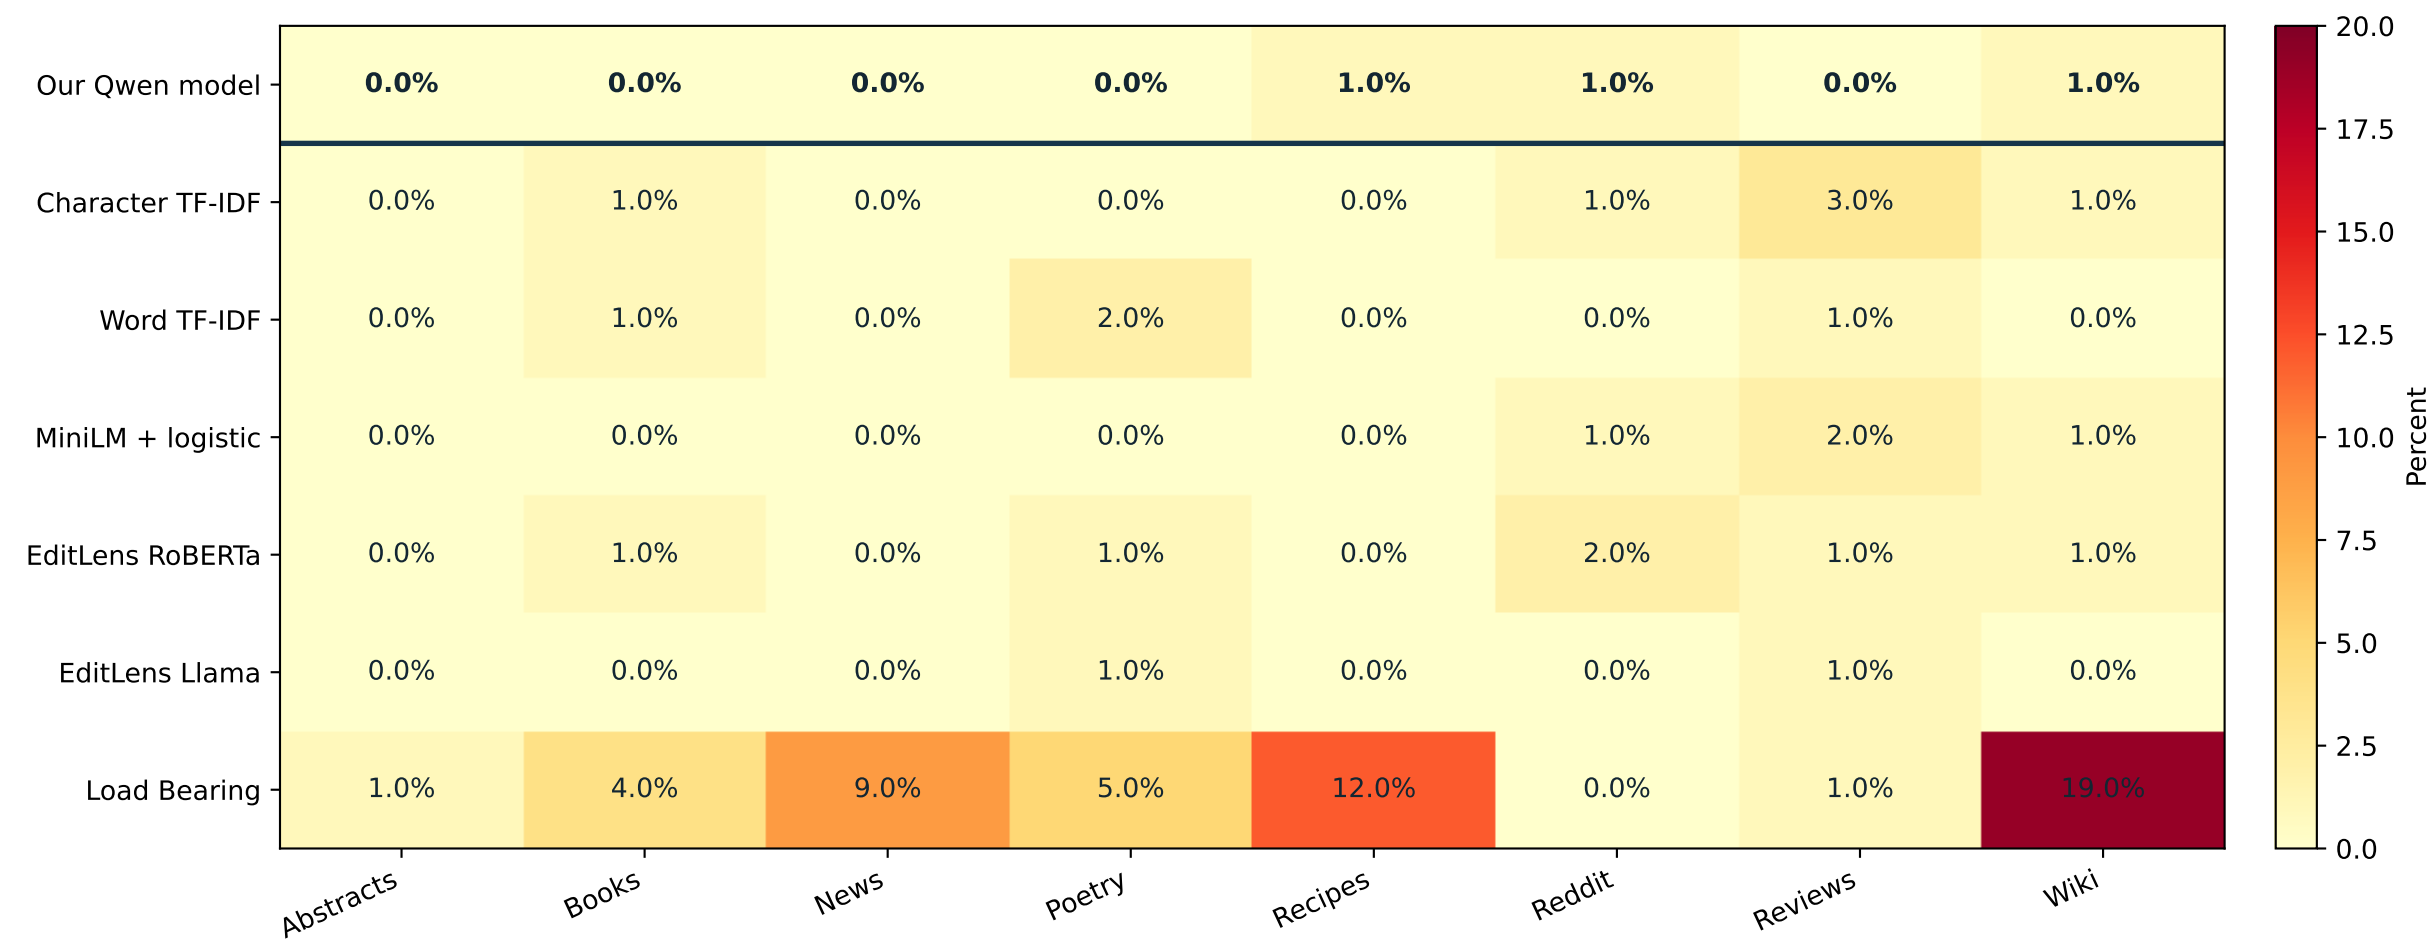


Passages per category (AI): Abstracts 100, Books 100, News 100, Poetry 100, Recipes 100, Reddit 100, Reviews 100, Wiki 100

Small categories have noisy percentages; one error changes a 50-passages category by 2 points.

RAID independent test: human false positives by category

Lower is better. Each cell is the percent of human passages wrongly flagged.

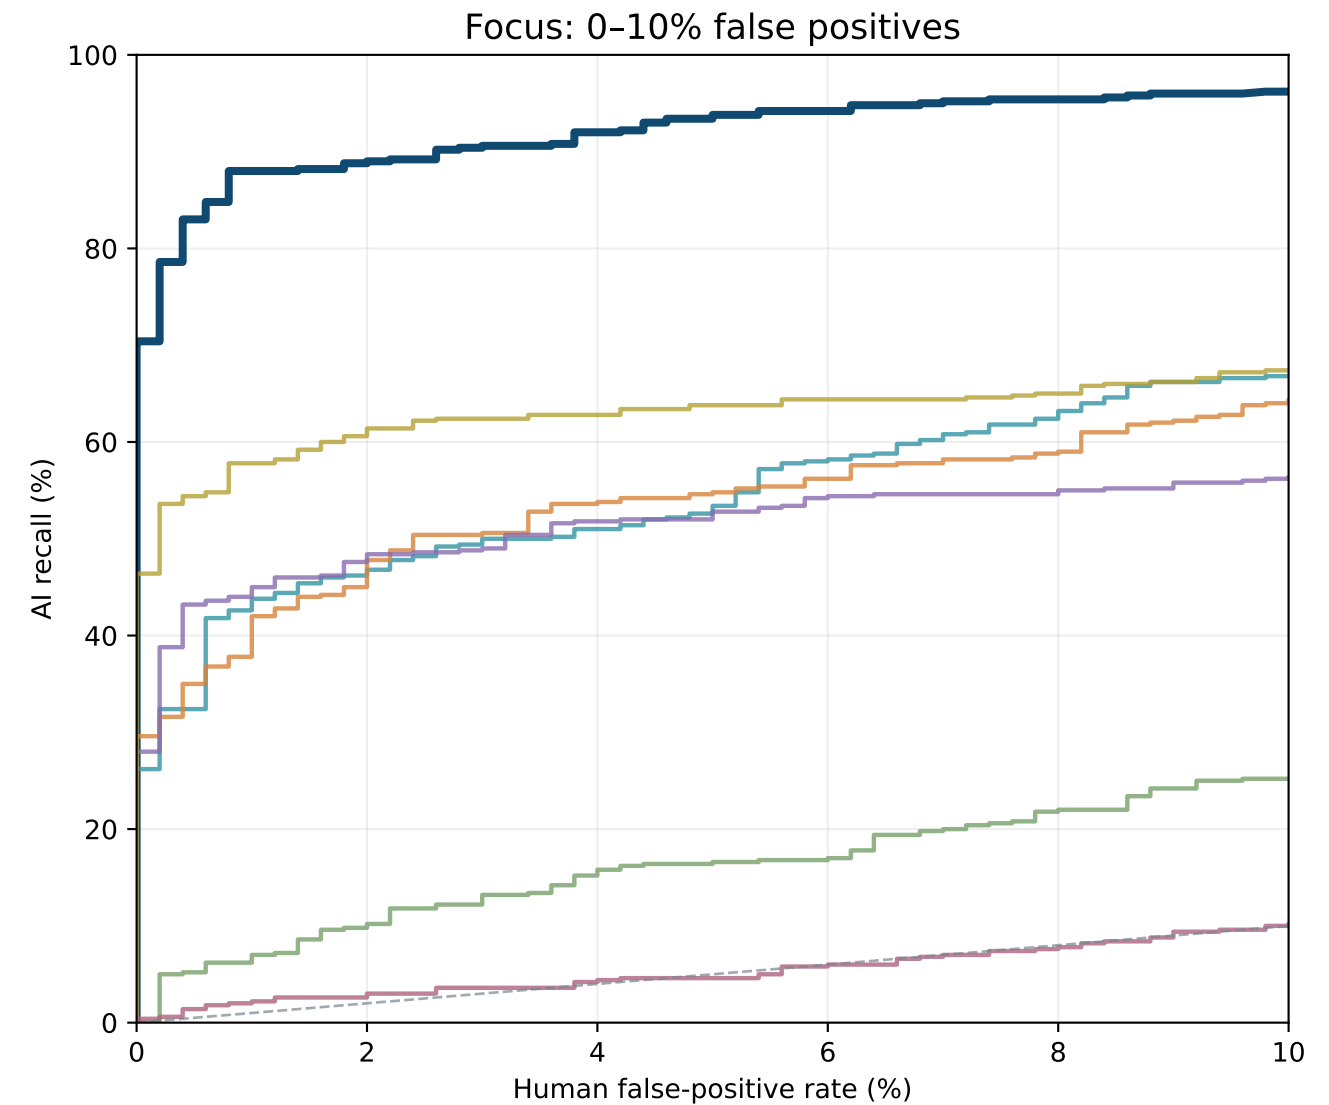
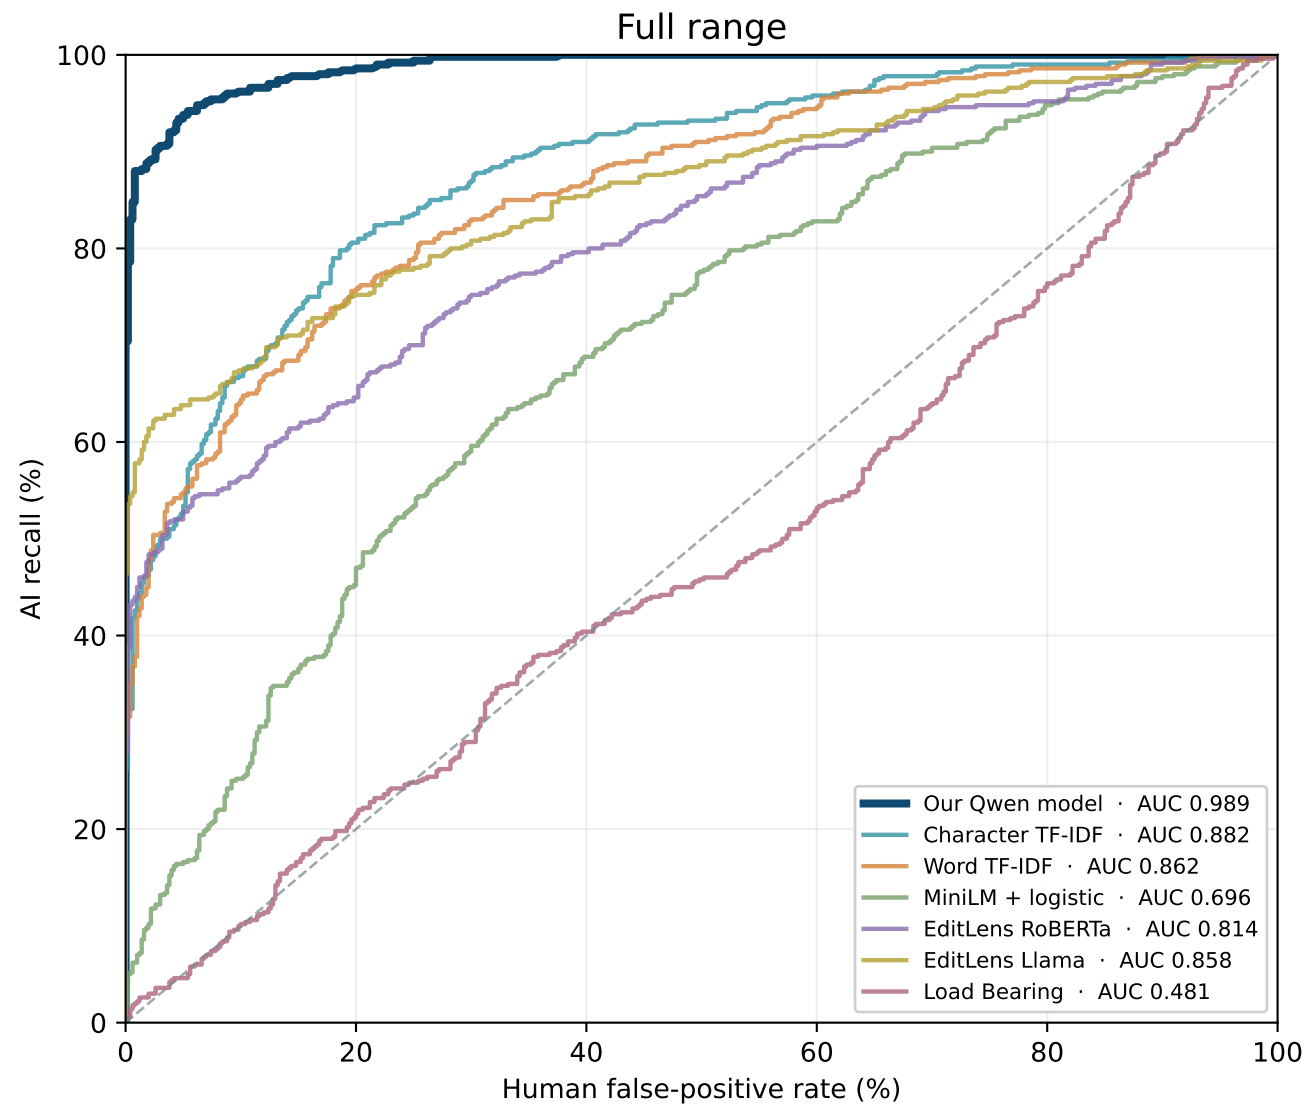


Passages per category (human): Abstracts 100, Books 100, News 100, Poetry 100, Recipes 100, Reddit 100, Reviews 100, Wiki 100

Small categories have noisy percentages; one error changes a 50-passages category by 2 points.

Diverse test: ROC curves

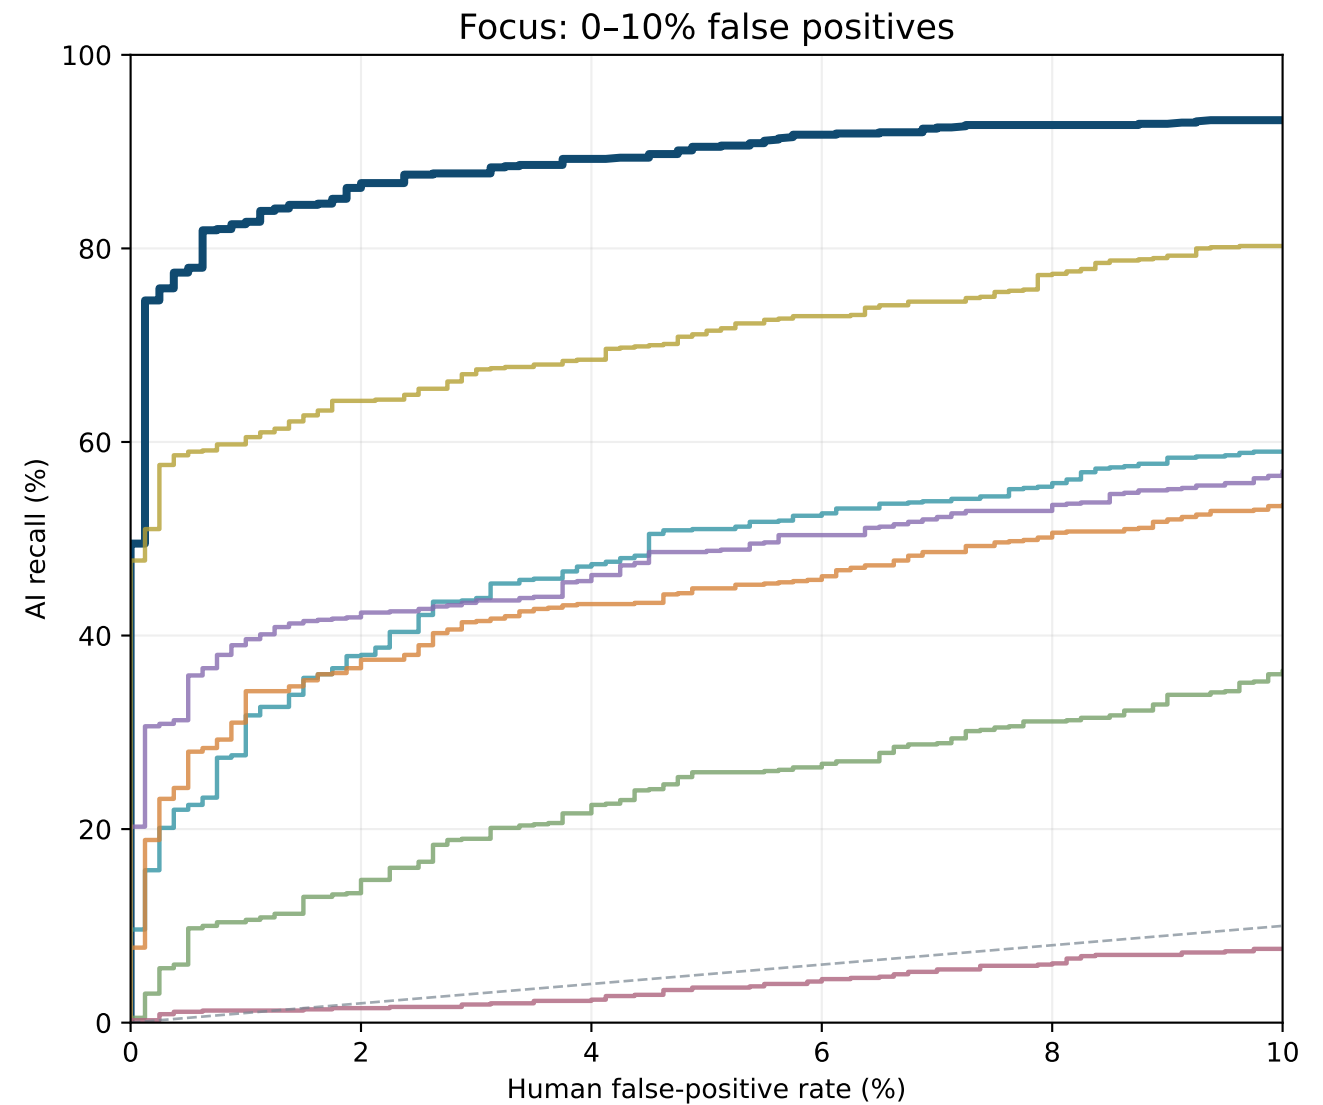
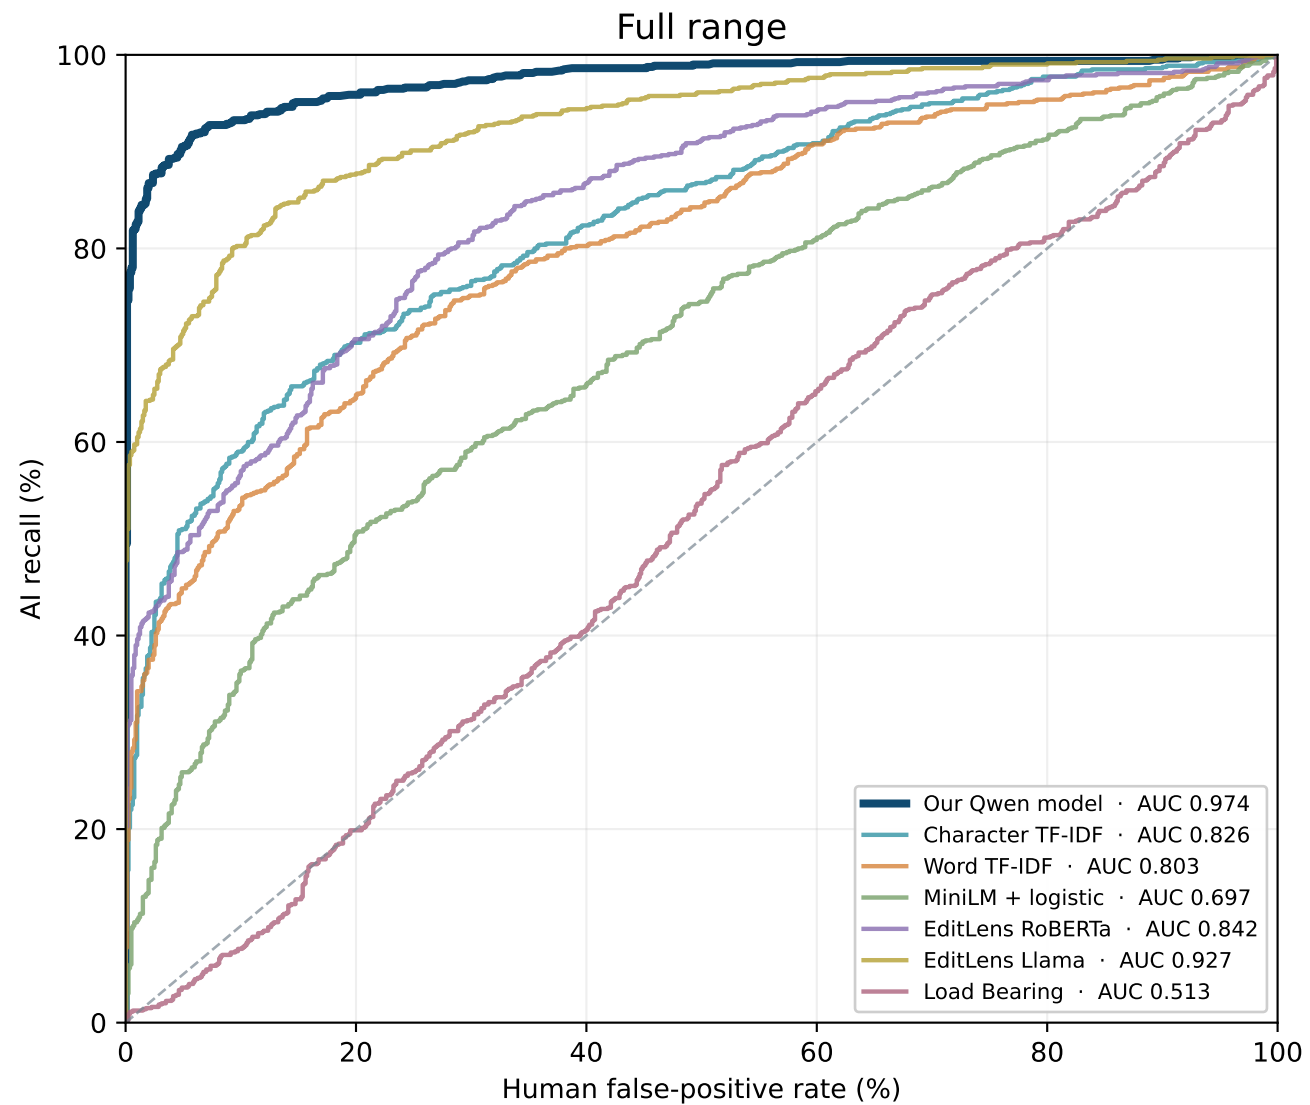
Each curve shows AI recall as the false-positive rate changes across all possible thresholds.



AUC summarizes ranking across all thresholds. It does not say which false-positive rate a deployed threshold will have.

RAID independent test: ROC curves

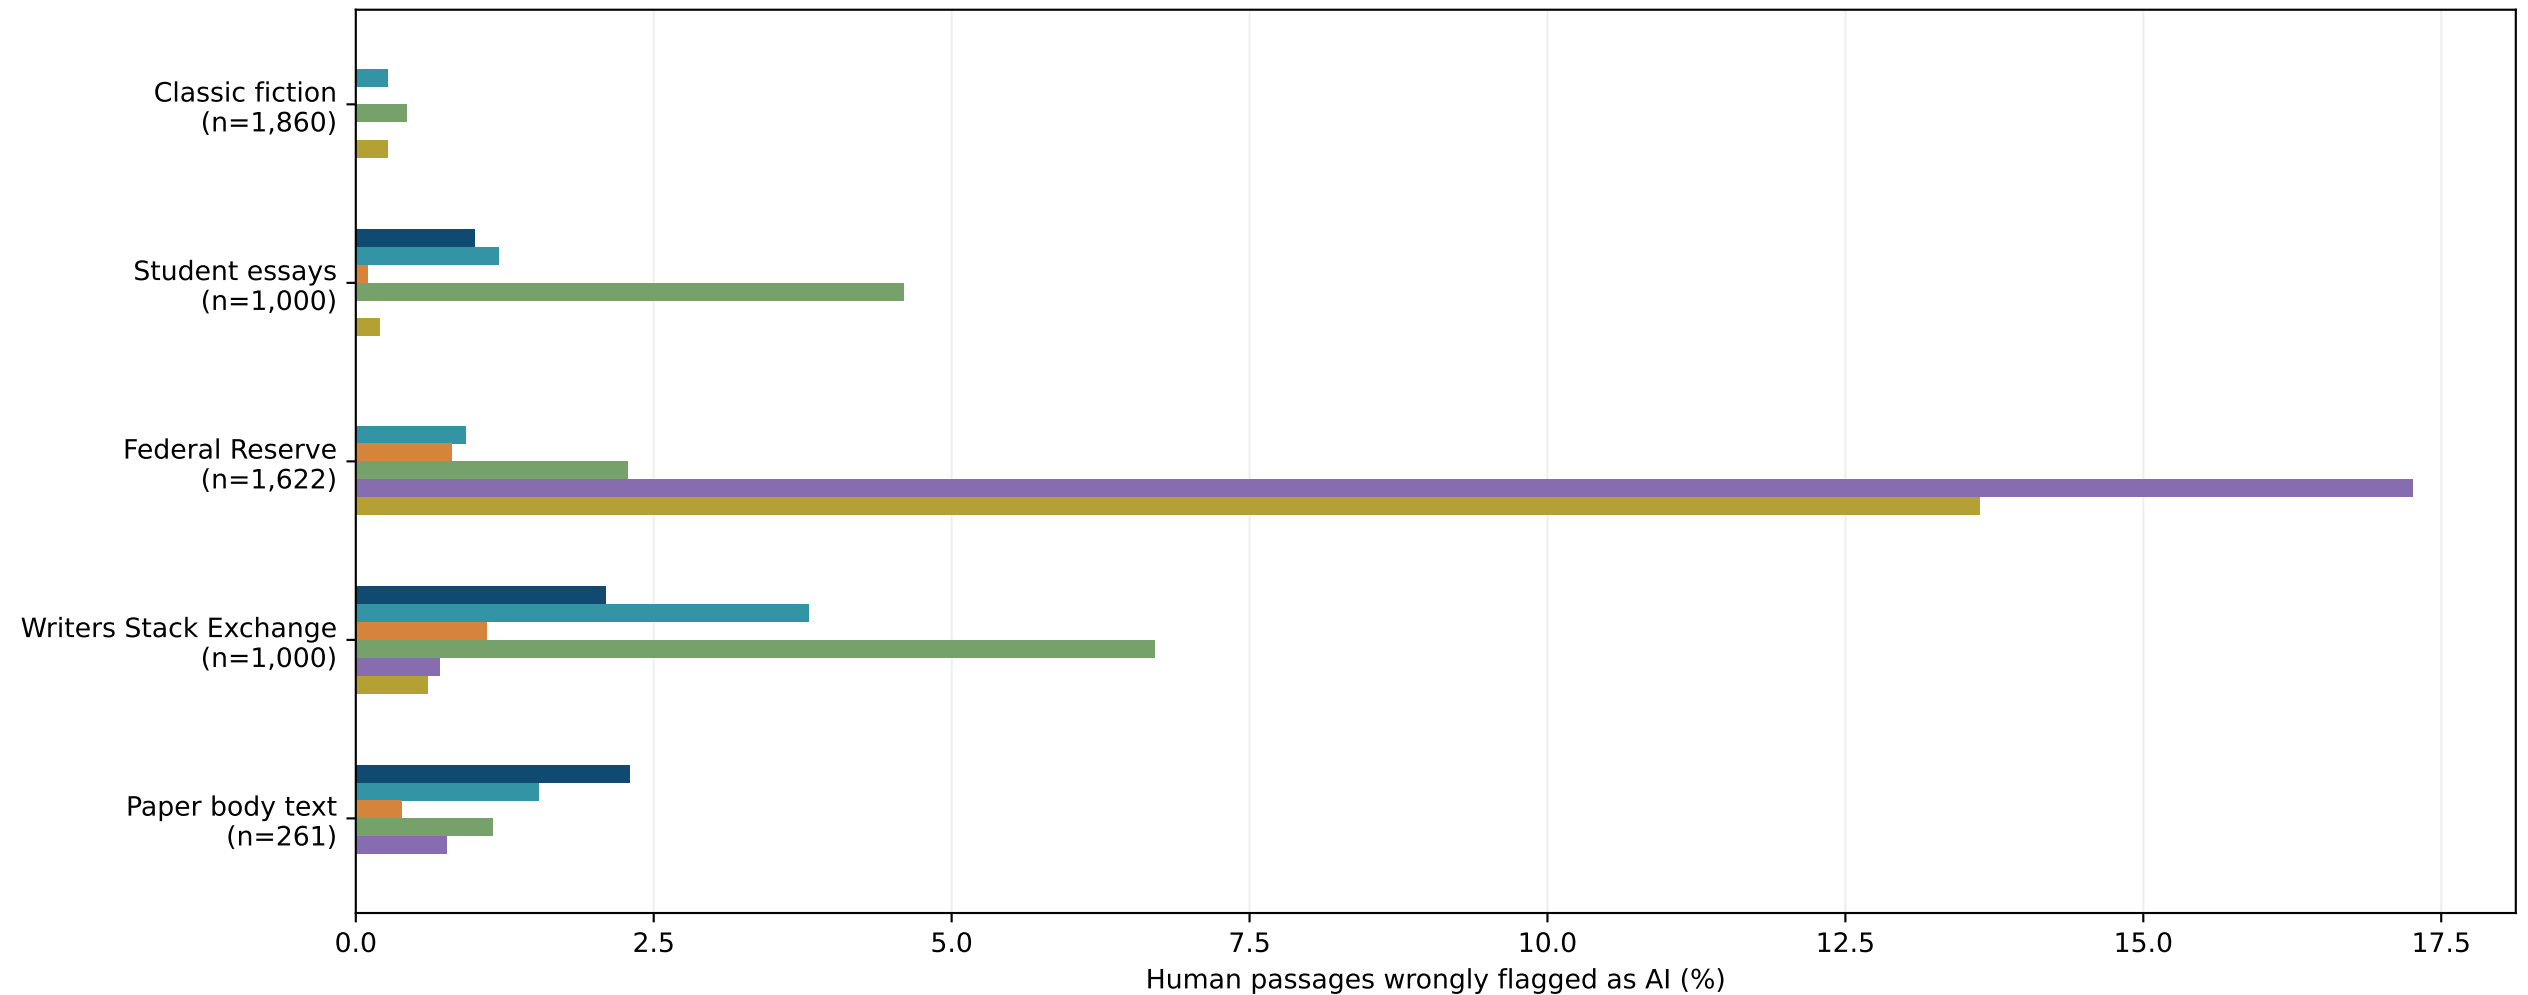
Each curve shows AI recall as the false-positive rate changes across all possible thresholds.



AUC summarizes ranking across all thresholds. It does not say which false-positive rate a deployed threshold will have.

Additional human-only checks

False positives on human text from sources outside the two balanced test sets. Lower is better.



These sets contain no AI examples, so they measure false positives only. Writers Stack Exchange authorship is inferred from pre-2023 posts.

