




External article stress test: frozen operating points



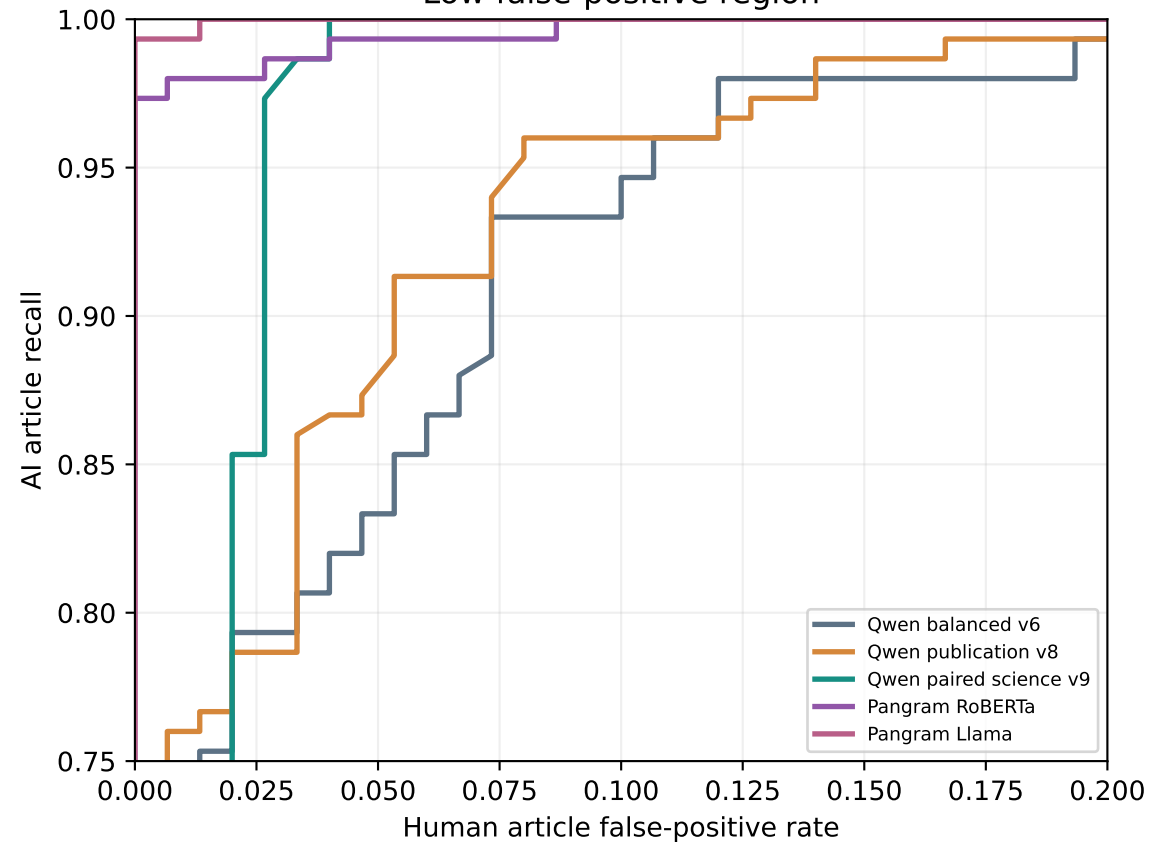
Dashed goals: $\leq 10\%$ human false alarms and $\geq 95\%$ AI article recall. Qwen flags an article if any token is highlighted; Pangram baselines make whole-article decisions.

External article discrimination across decision thresholds

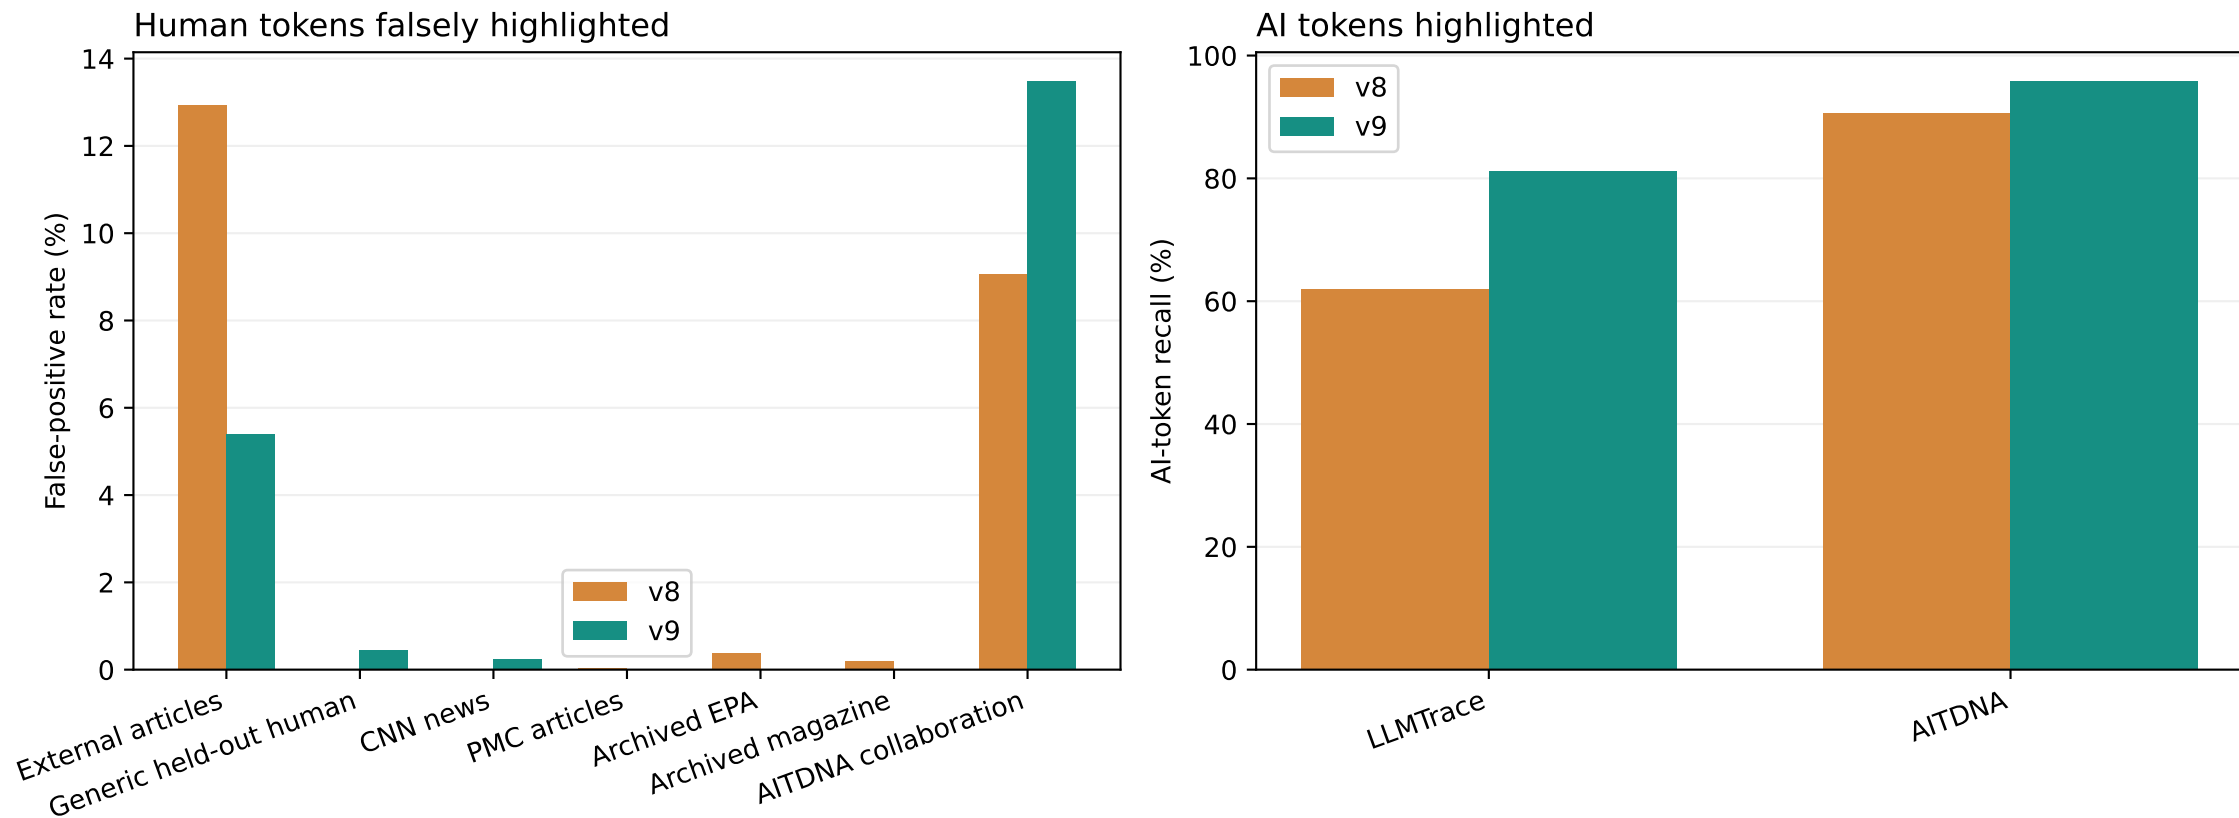
Full ROC



Low false-positive region

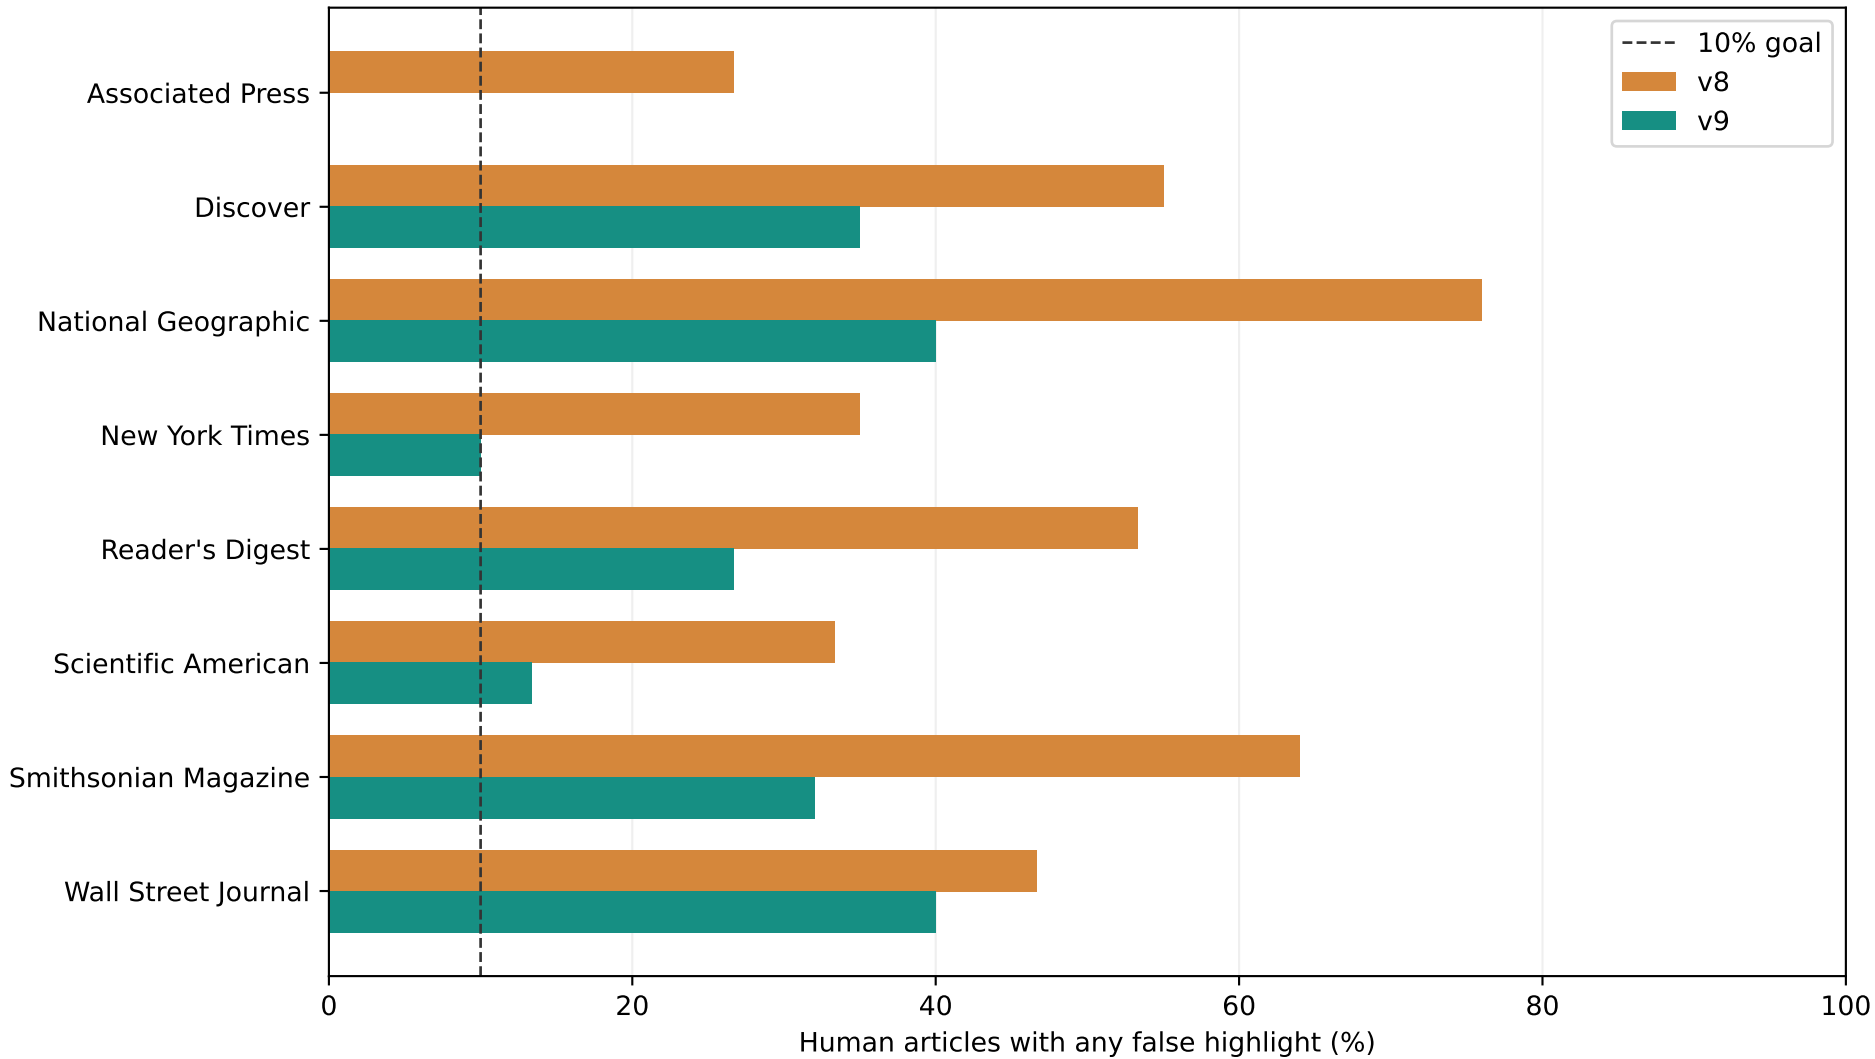


Frozen-threshold transfer across domains



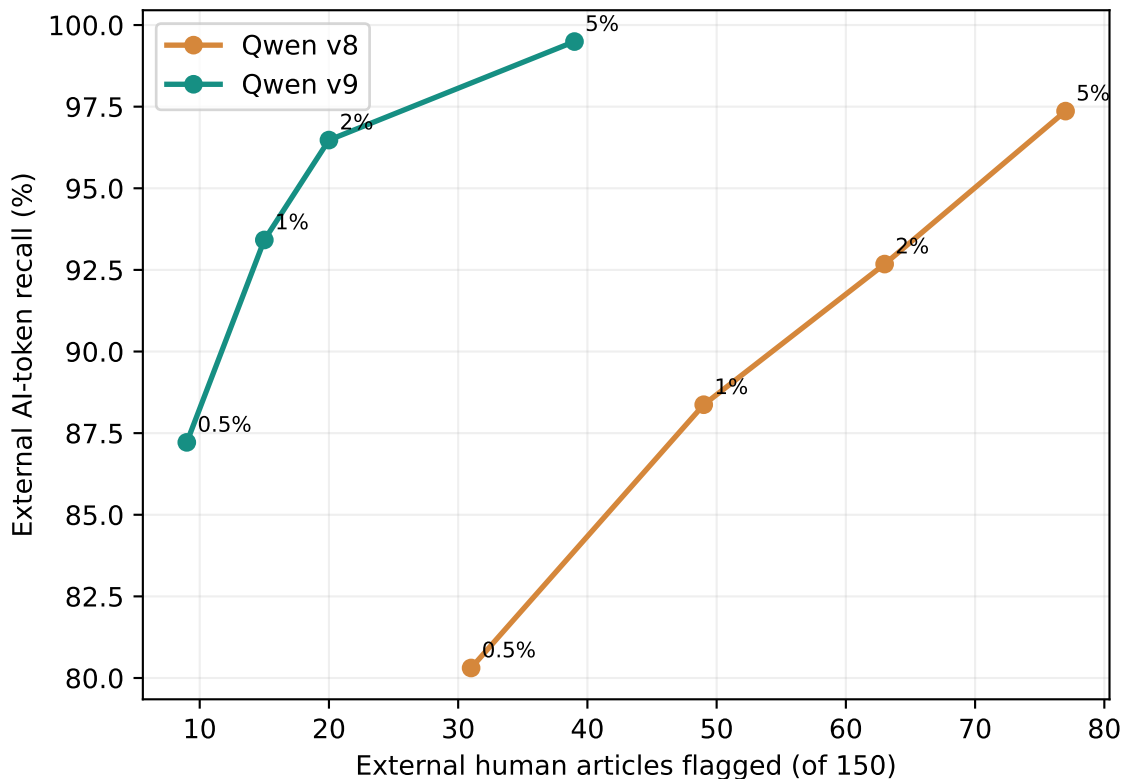
Archived magazine test authors were excluded from its training candidates. AITDNA contains real human-AI collaboration with token-level provenance.

False alarms by publisher on external articles

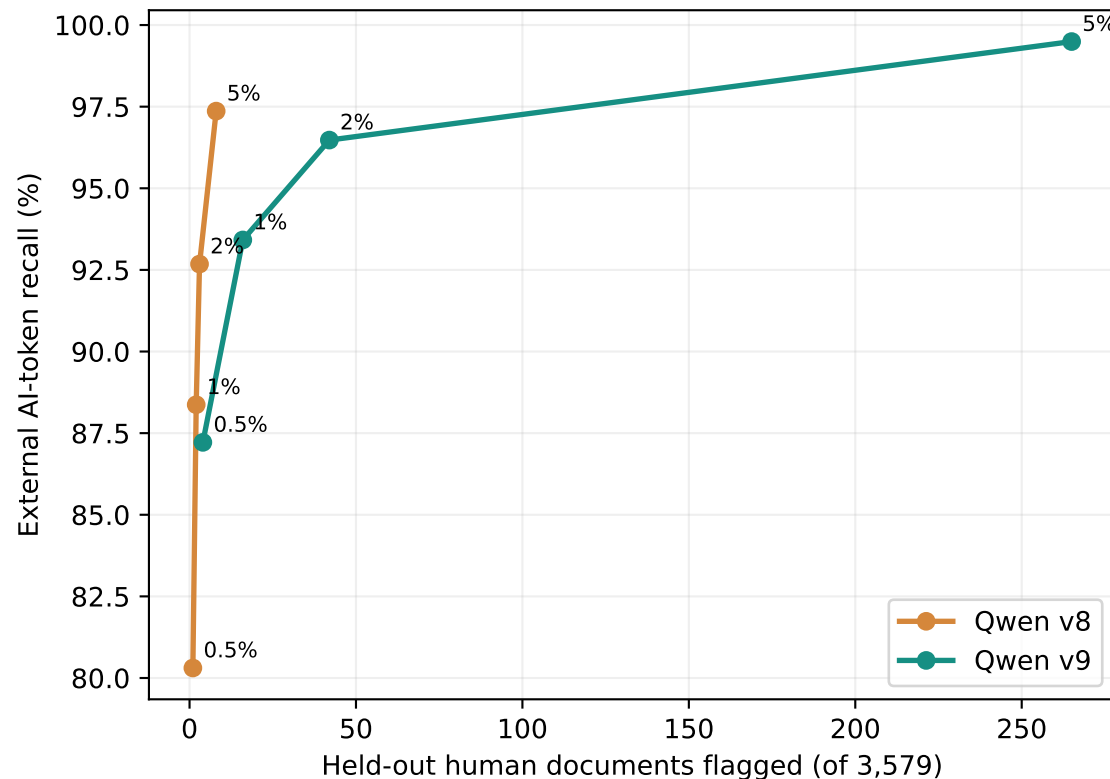


Thresholds selected on separate human calibration only

Publication transfer



General human transfer



Point labels are target false-alarm rates on the calibration documents. No held-out examples were used to set these thresholds.