



Held-out LLMTrace: 2,000 documents / nine domains

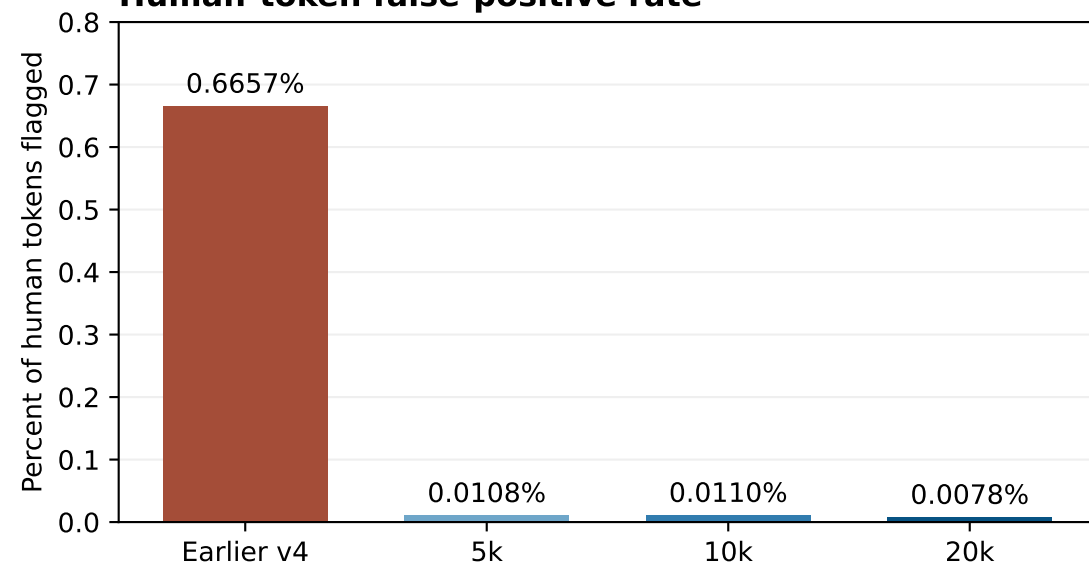
AI-token recall, all documents



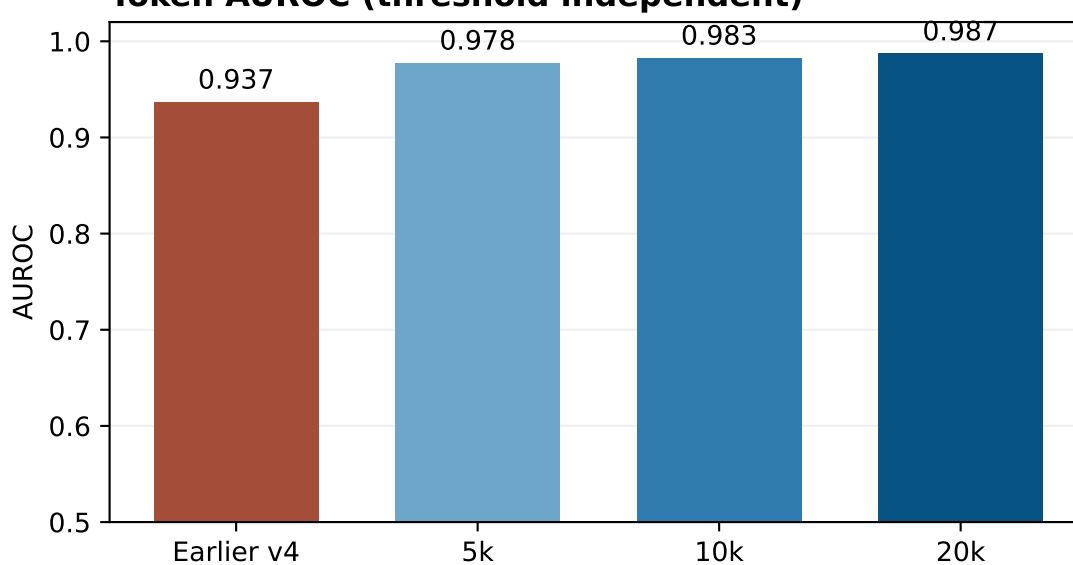
AI-token recall, mixed documents only



Human-token false-positive rate

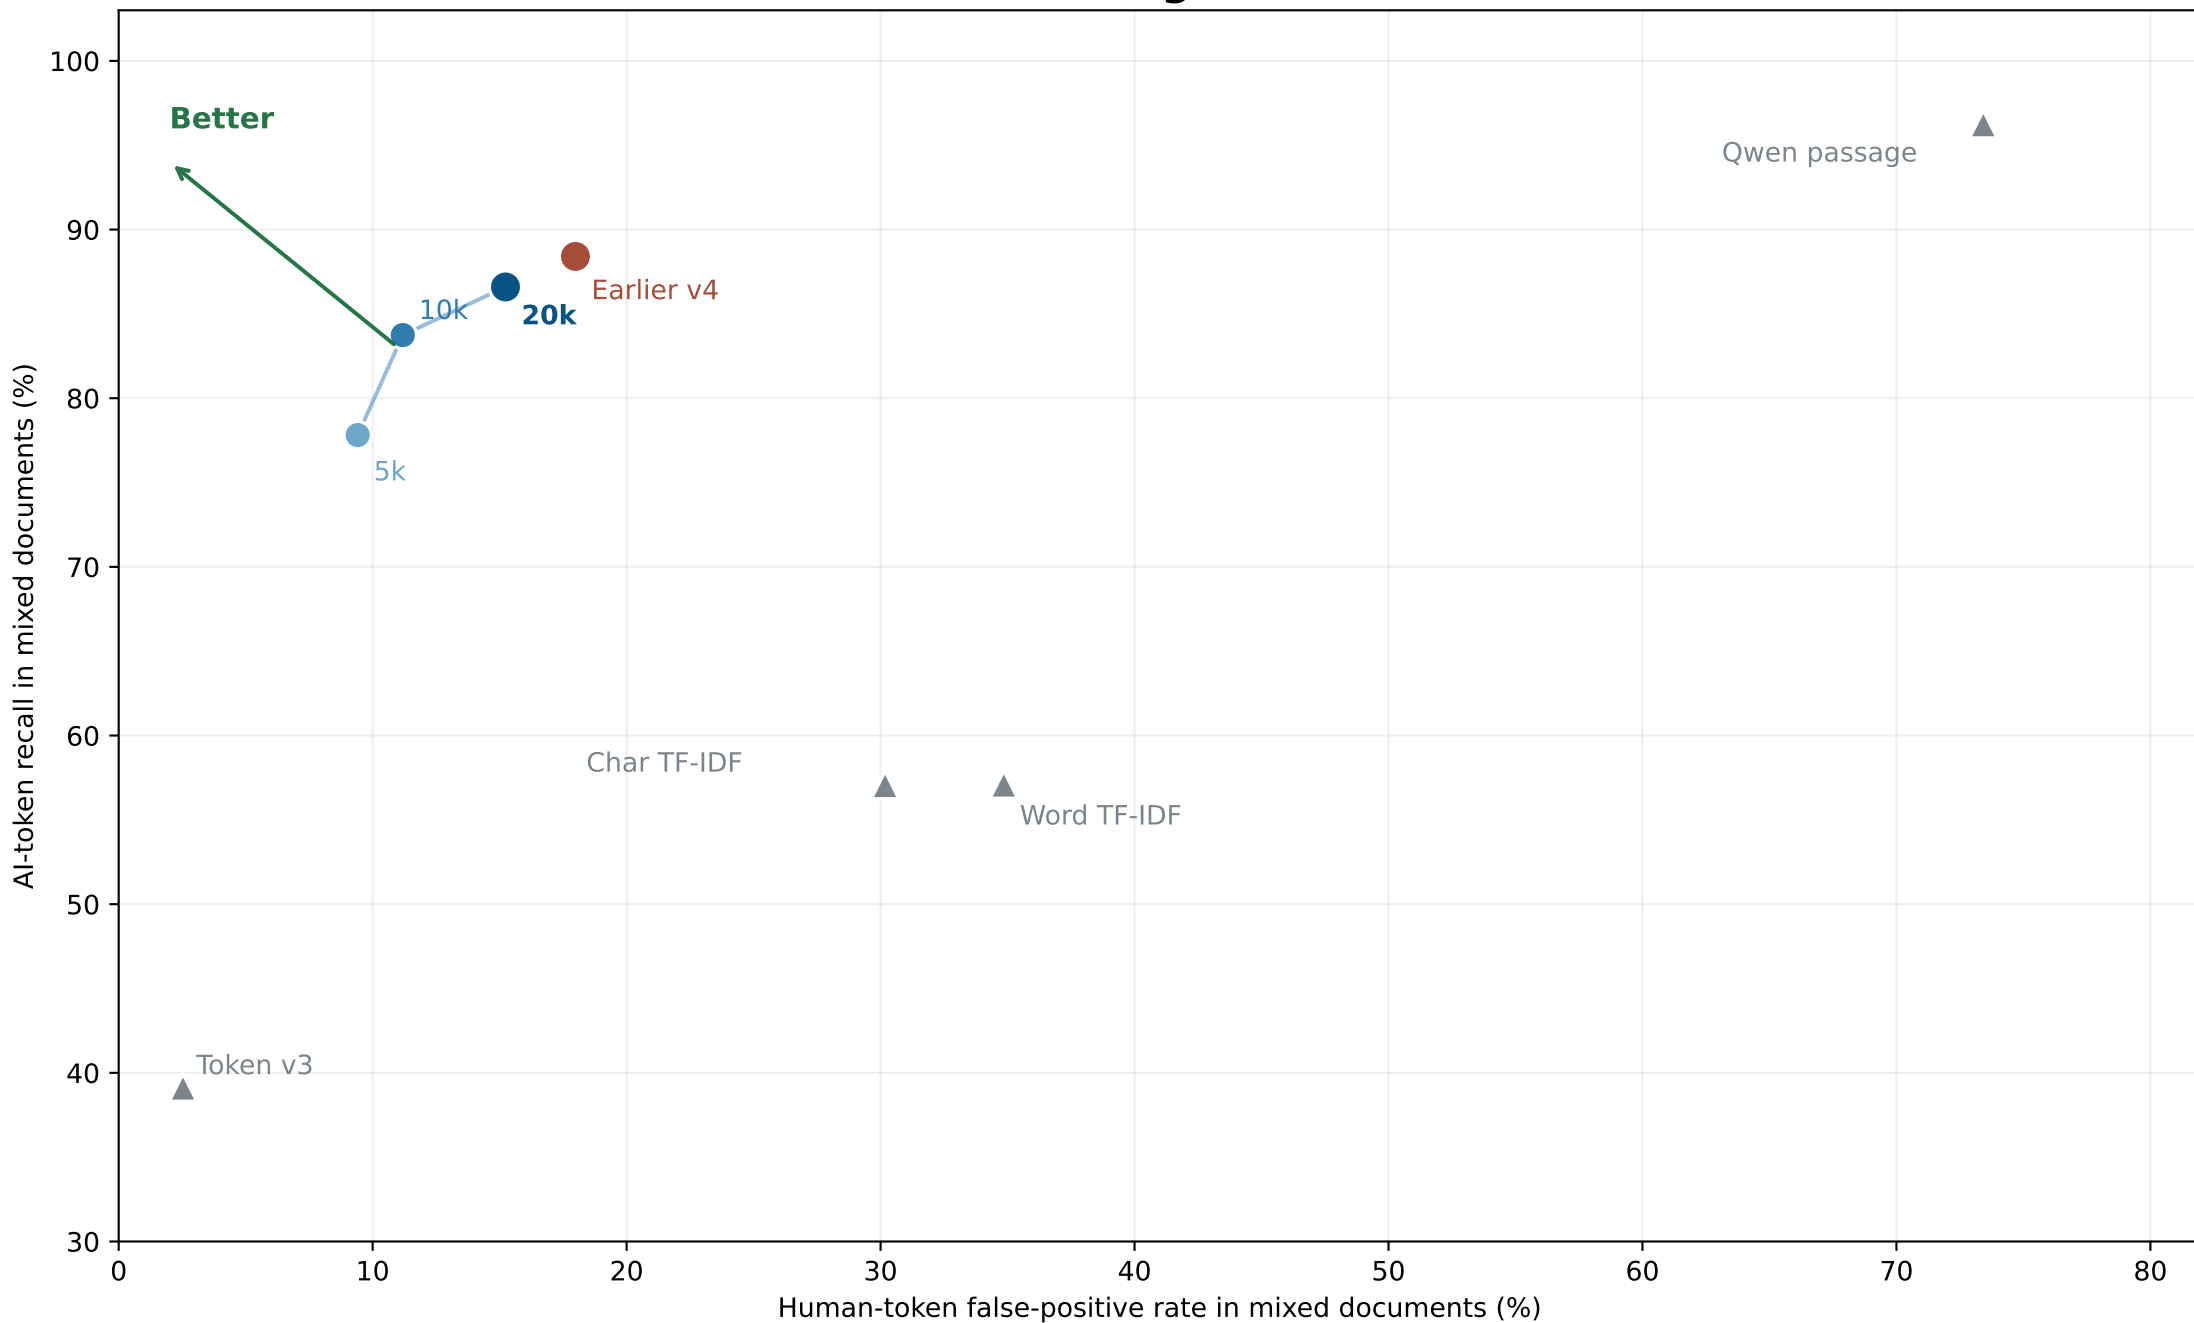


Token AUROC (threshold independent)



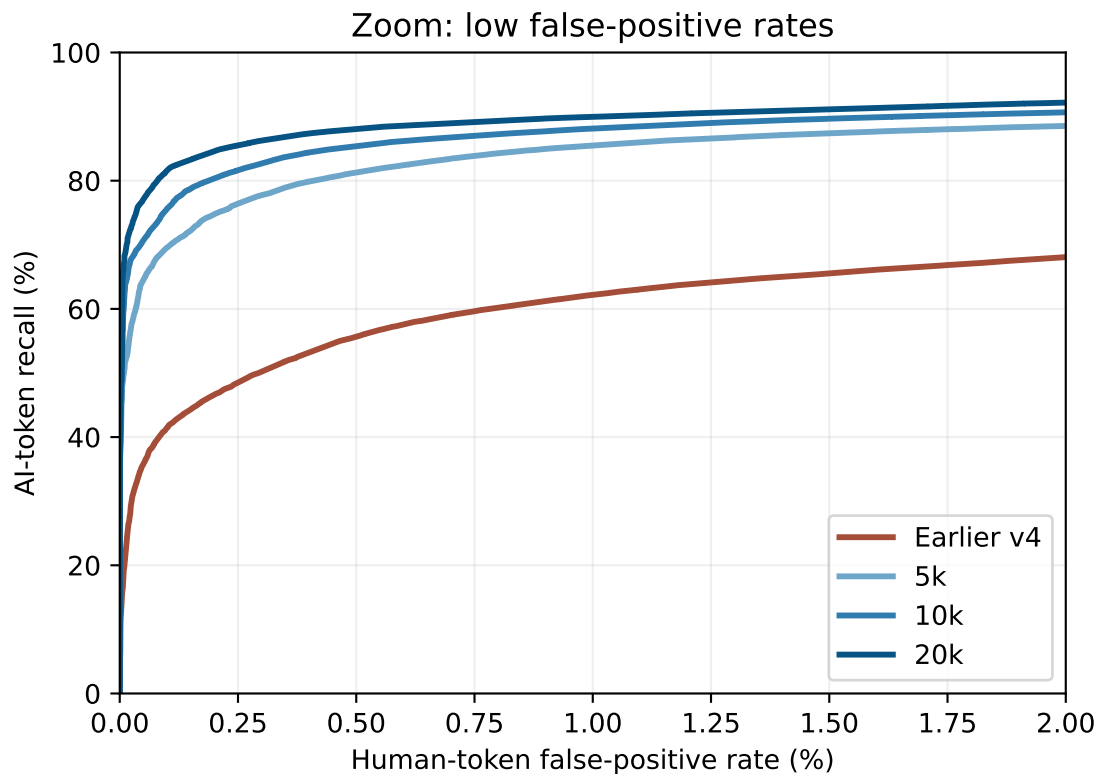
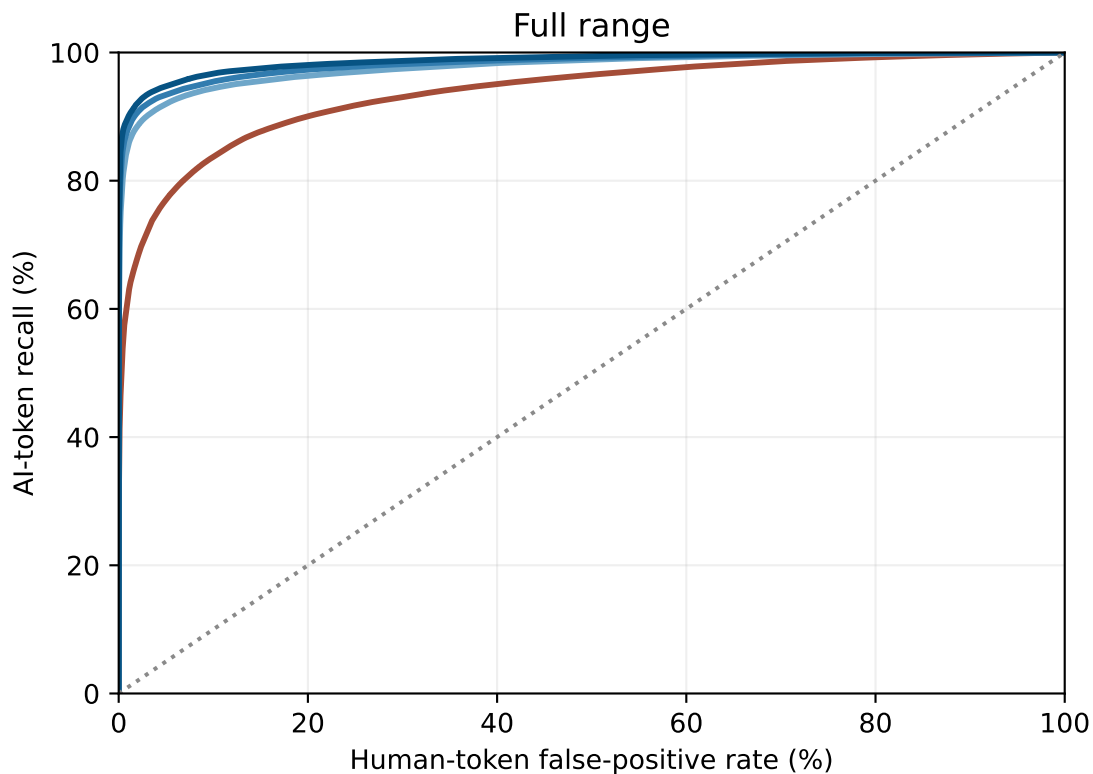
Frozen threshold per model: calibrated on the same separate 1,120-document human set.
Earlier v4 used a different 5k training mixture; 5k/10k/20k are nested and comparable.

AITDNA: real human-AI collaborative writing (258 mixed documents)



Each model uses its own threshold calibrated on the same separate pure-human set.
At these thresholds, CoAuthor AI-token recall is 0.0% (5k), 0.6% (10k), and 0.0% (20k).

LLMTrace ROC: available recall / false-positive tradeoff



These curves use test-set scores to show possible tradeoffs; threshold choices on this plot are diagnostic, not deployable.